A Practical Introduction to Data Science

Part 1

Gergely Zsombor Haász

haasz.zsombi@gmail.com

Course Agenda

I.	Introduction to Data Science
II.	Business and Data Understanding
III.	Introduction to Supervised Learning
IV.	Advanced Supervised Learning
V.	Unsupervised Learning
VI.	Time Series Analysis
VII.	Deep Learning
VIII.	Machine Learning Operations

Introduction to Data Science

- Data Science
- Machine Learning
- The ML Lifecycle
- ML Algorithms
- The Goal
- Deep Learning
- Tools

Data Science

The Data Analytics Iceberg

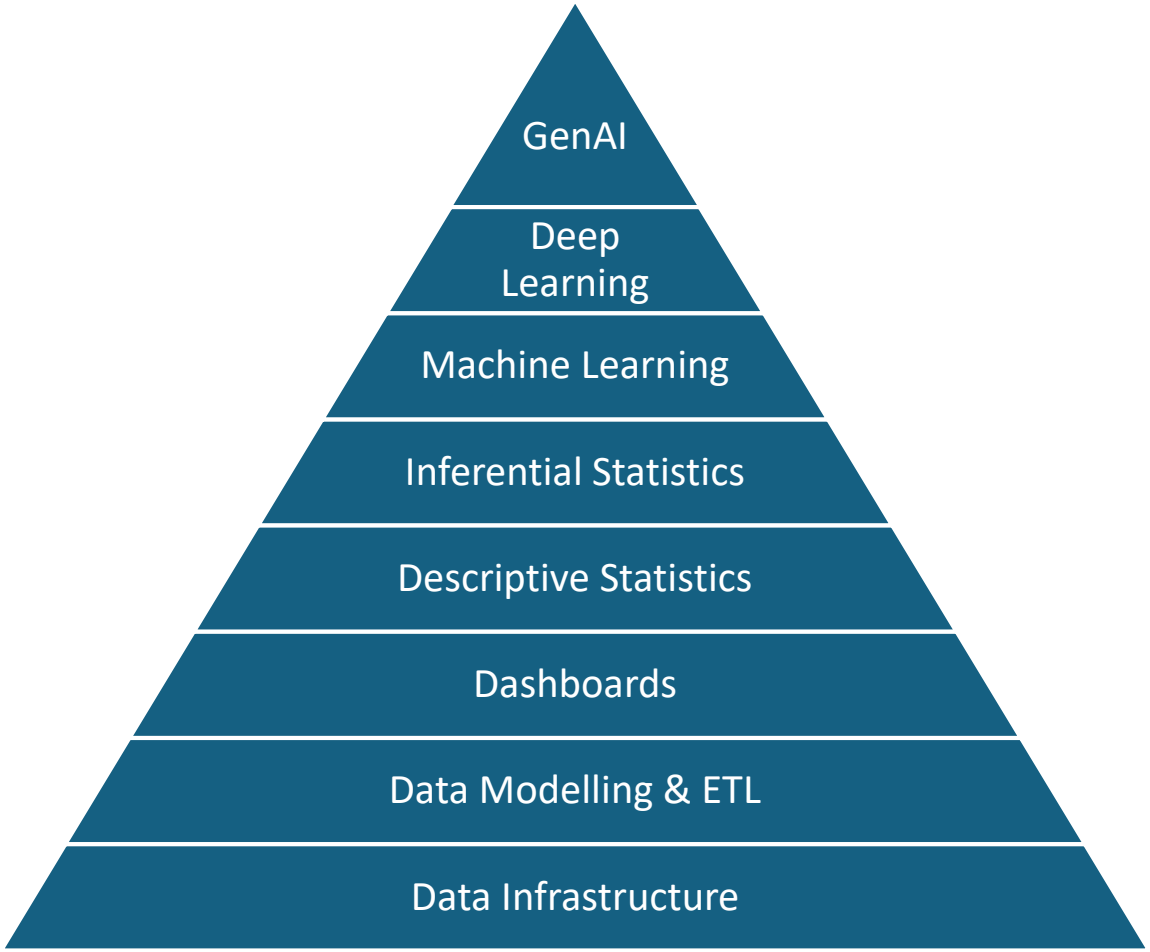

The Data Analytics Iceberg

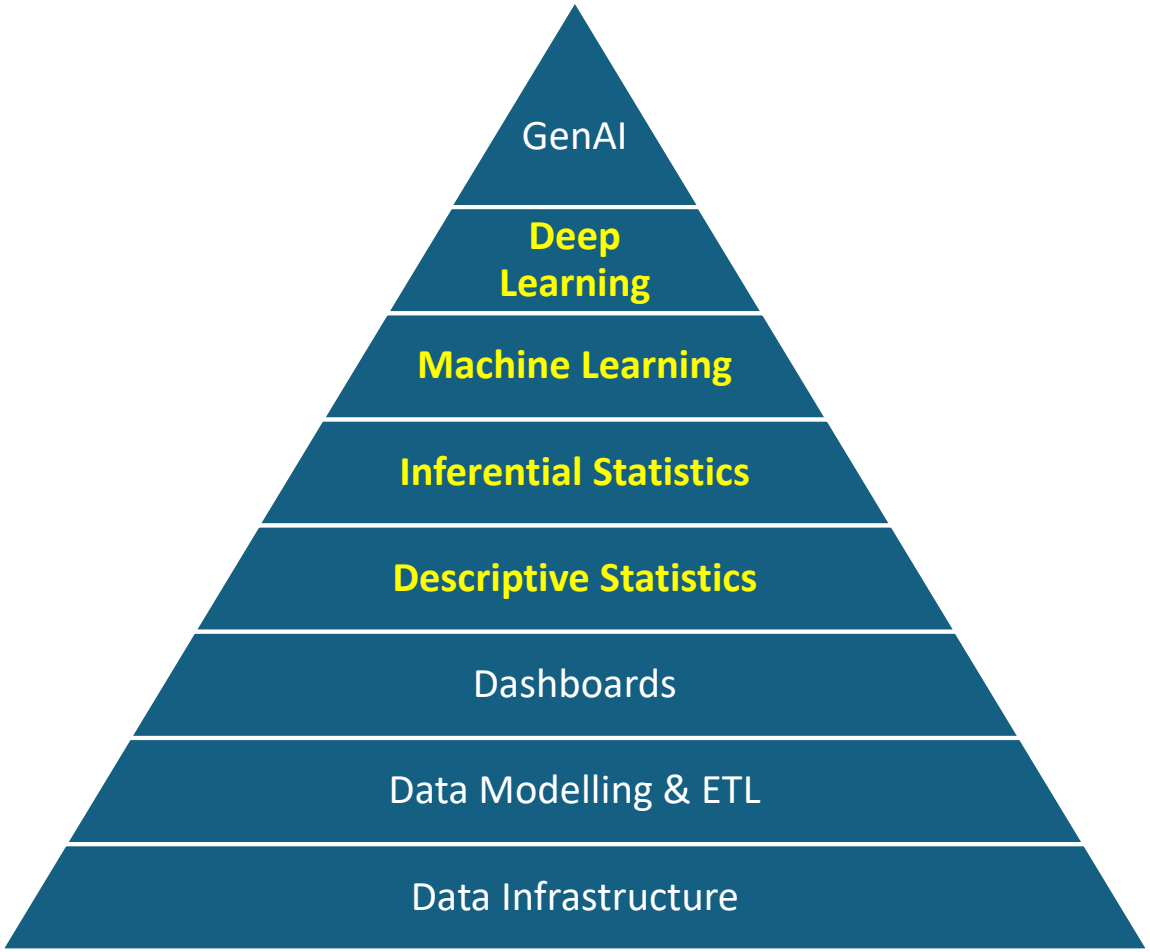

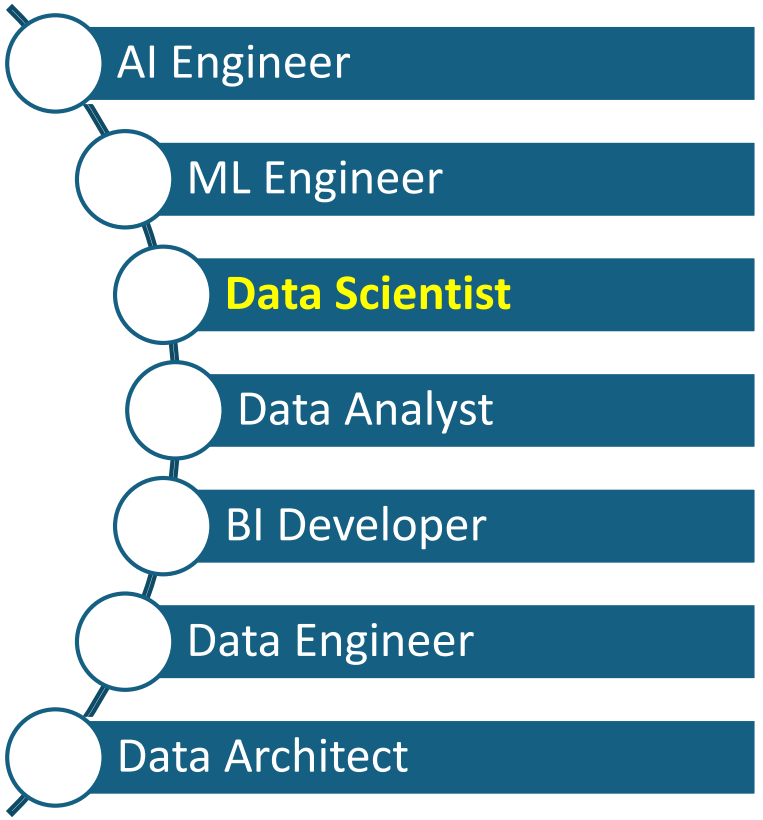

What is Data Science?

Data Science *is about* ***extracting information*** *from data by* ***finding patterns****, to make better data-driven decisions.*

Data types	Methods	Applications
• Tabular • Image & video • Text & speech • Other unstructured	• Descriptive statistics • Dashboards • Inferential statistics • Machine Learning • Deep Learning • Optimization	•healthcare, psychology, finance, marketing, entertainment, sport, transportation, logistic, manufacture, energy, telecommunications, customer relationships, •etc.

Data & Science

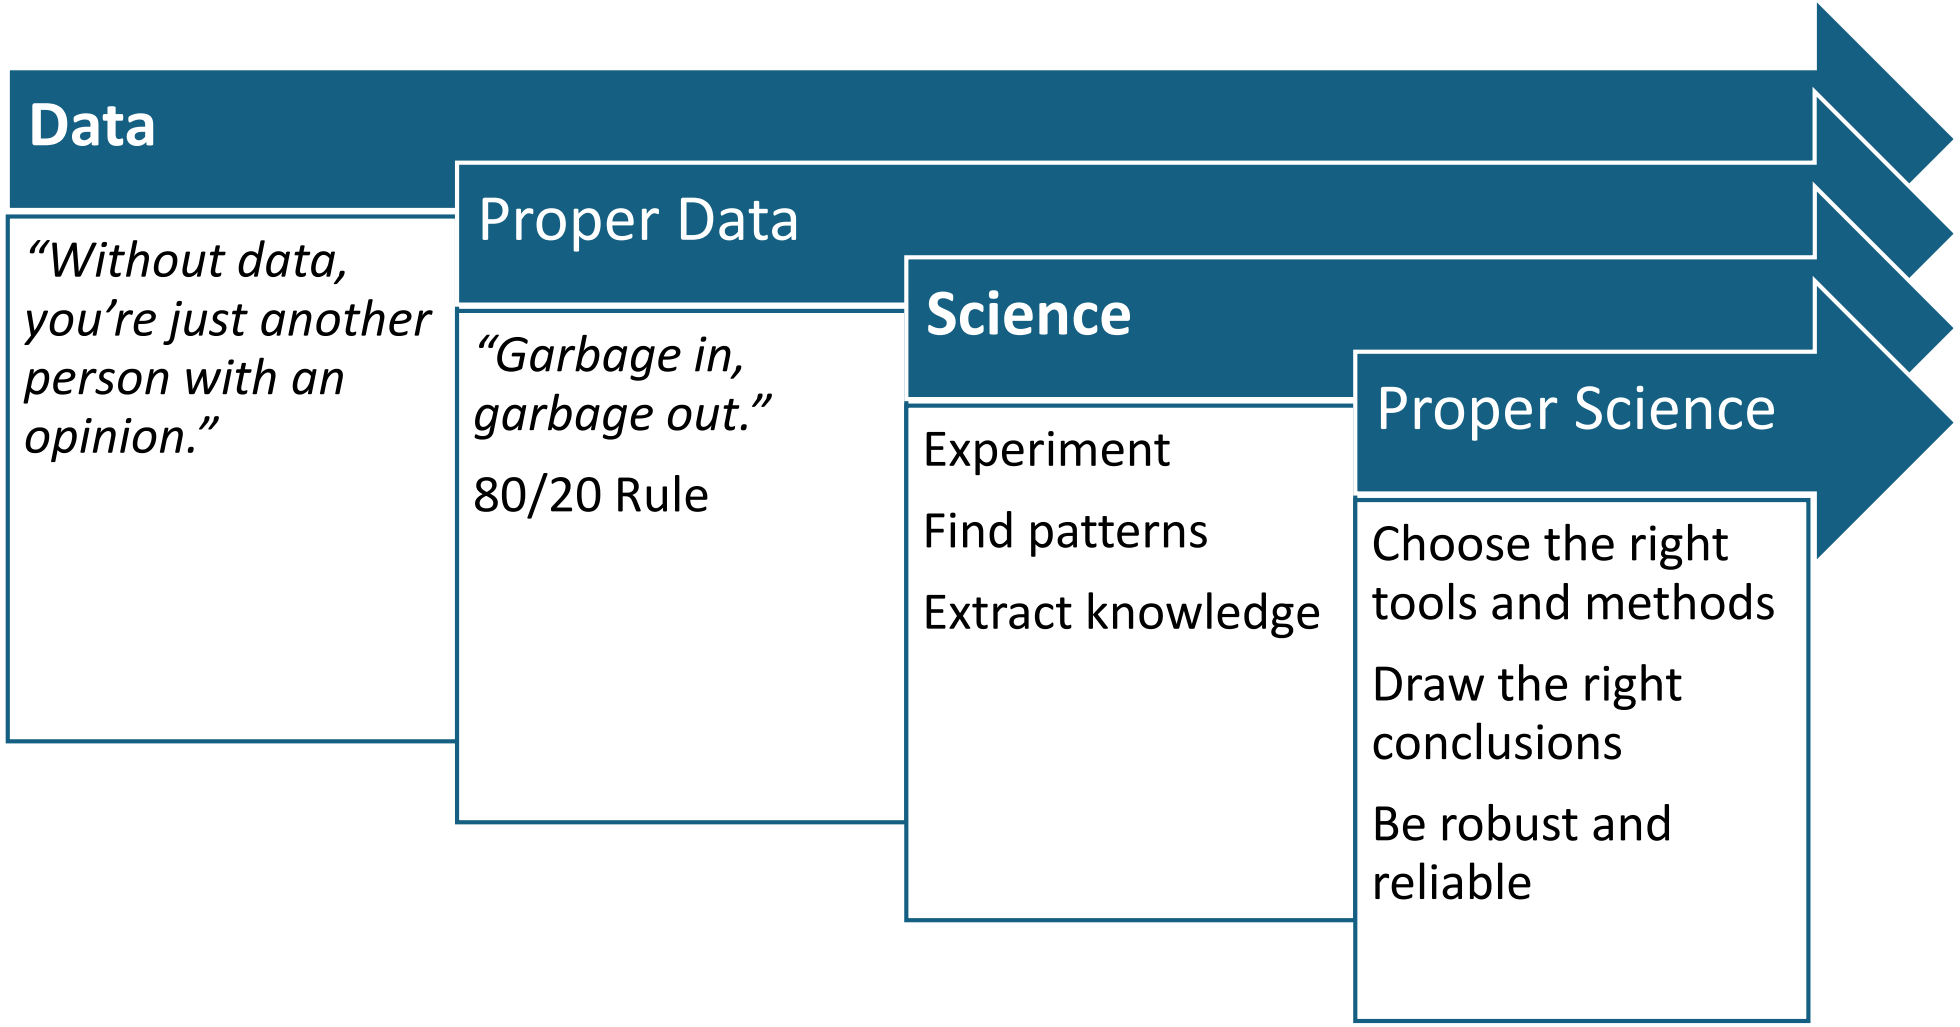

Machine Learning

Statistics

Descriptive

- Goals:
  - Understand as it is
  - Gather insights
- Tools:
  - Measures of Central Tendency
  - Measures of Dispersion
  - Distribution Properties
  - Correlation

Inferential

- Goals:
  - Generalization
  - Measure uncertainty
  - Causality
- Tools:
  - Estimation
  - Hypothesis Testing
  - Regression Analysis

What is Machine Learning?

”the ability to learn without being explicitly programmed” (Arthur Samuel, 1959)

An **ML model** is simply a mathematical function or algorithm that maps **input data** to an output (**prediction**).

Instead of explicit programming, an **optimization algorithm** is used to determine the best **parameters** of the model that minimize the prediction **error**.

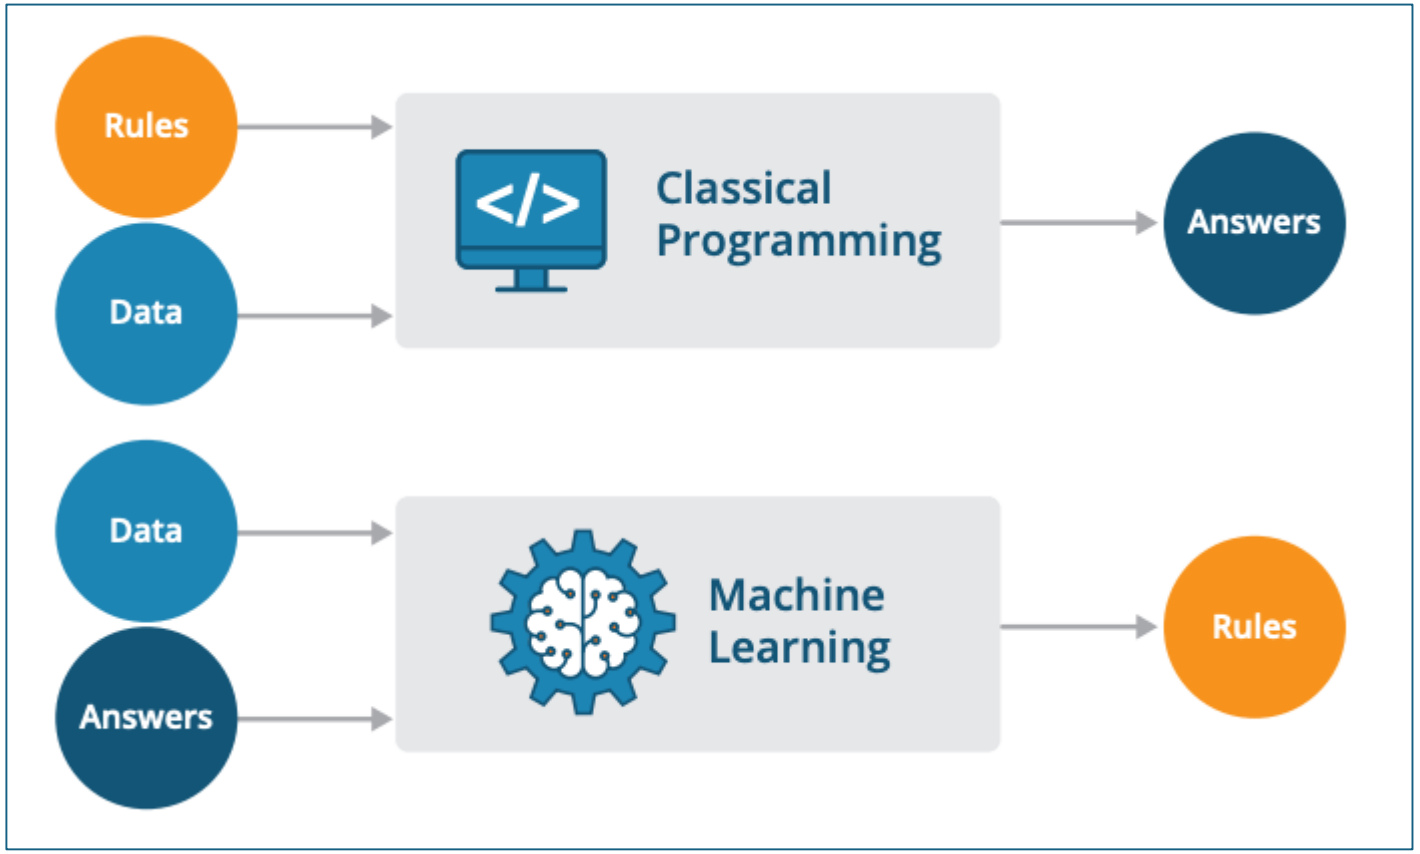

Source: Machine Learning is Forcing Software Development to Evolve - ML in Production

What is Machine Learning?

”the ability to learn without being explicitly programmed” (Arthur Samuel, 1959)

When to use Machine Learning?

1. Large data
2. Automation
3. Complex or changing patterns

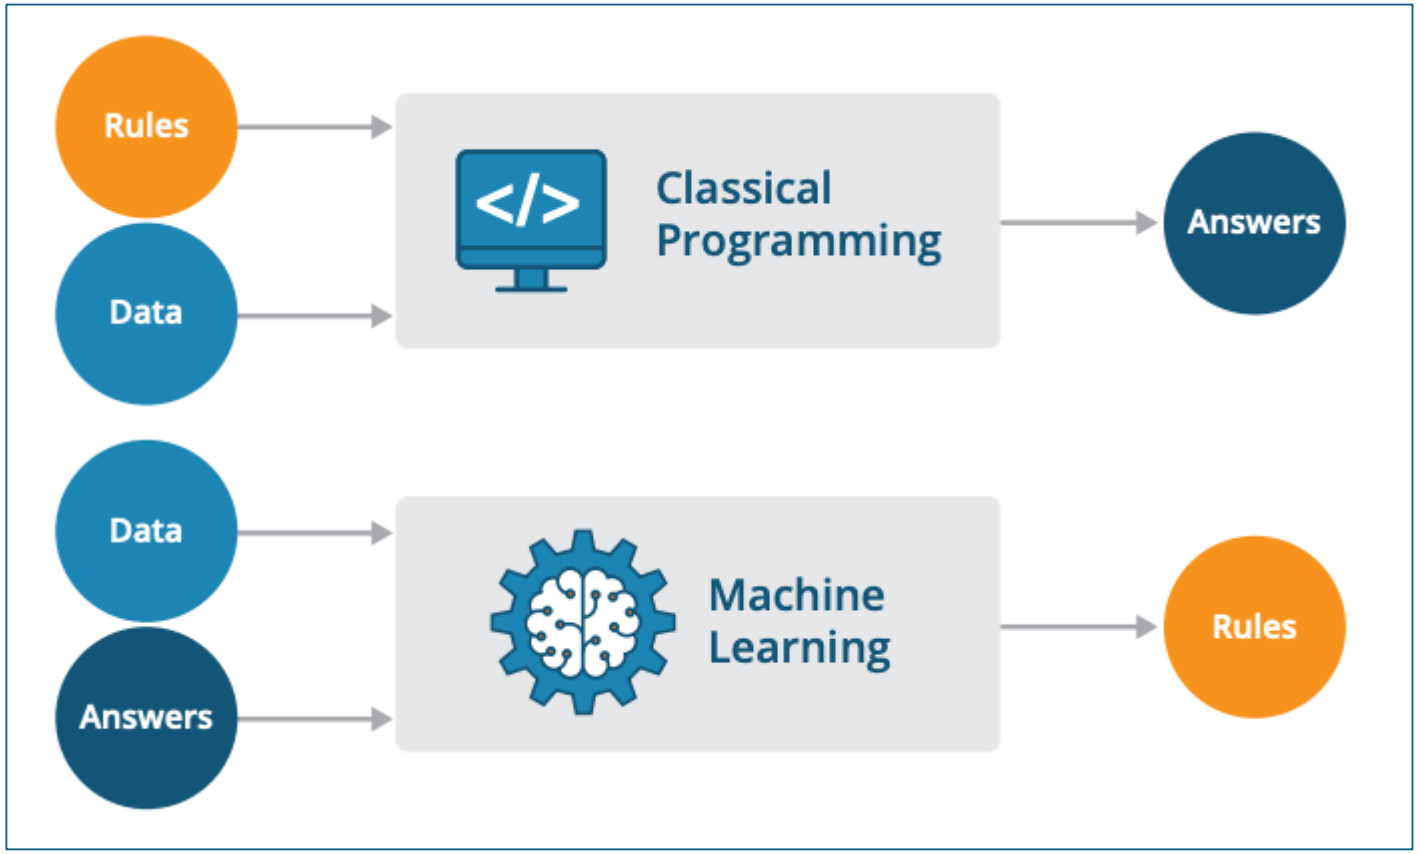

Source: Machine Learning is Forcing Software Development to Evolve - ML in Production

What is Machine Learning?

Supervised Learning

predict, forecast, explain

- ☐ Classification
- ☐ Regression
- ☐ Survival Analysis

Unsupervised Learning

recognize patterns & structure

- ☐ Clustering
- ☐ Dimensionality Reduction
- ☐ Anomaly Detection
- ☐ Recommendation Systems

Reinforcement Learning

optimize actions

- ☐ Model-based
- ☐ Model-free

Supervised Learning

Supervised Learning

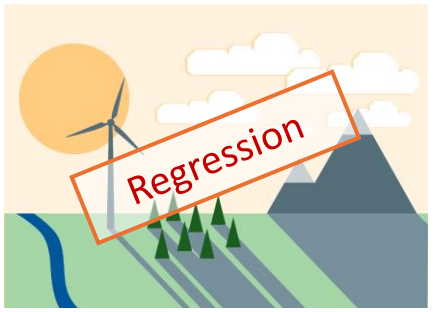

Unsupervised Learning

Clustering

- Customer segmentation

Dimensionality Reduction

- Noise reduction, Visualization, Latent Variables

Anomaly Detection

- Fraud detection, fault detection

Recommendation Systems

- Personalized product/movie/news recommendations

Reinforcement Learning

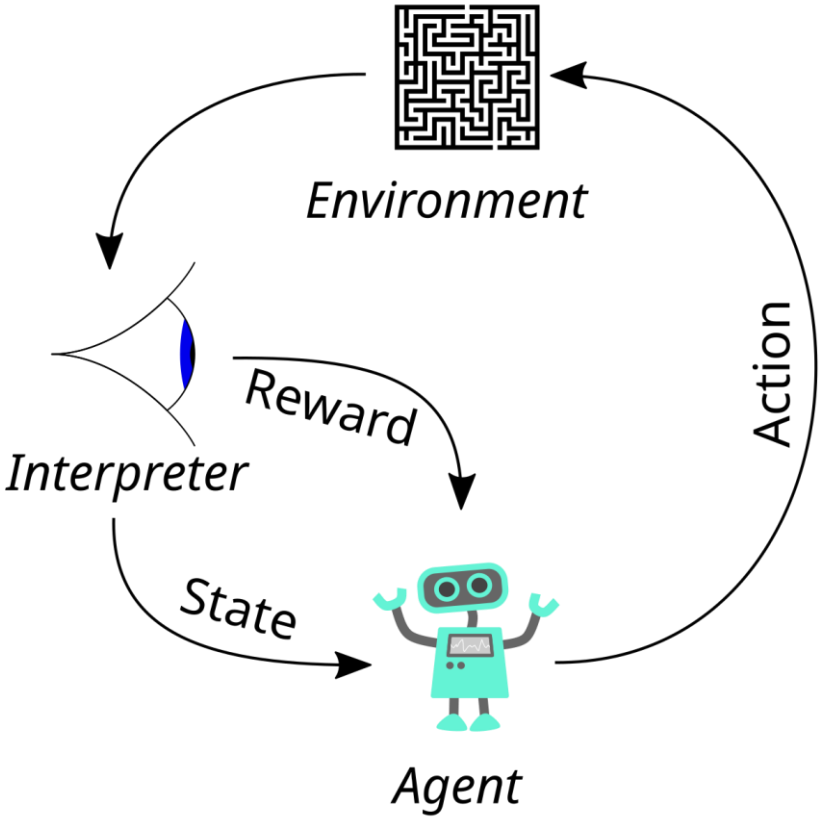

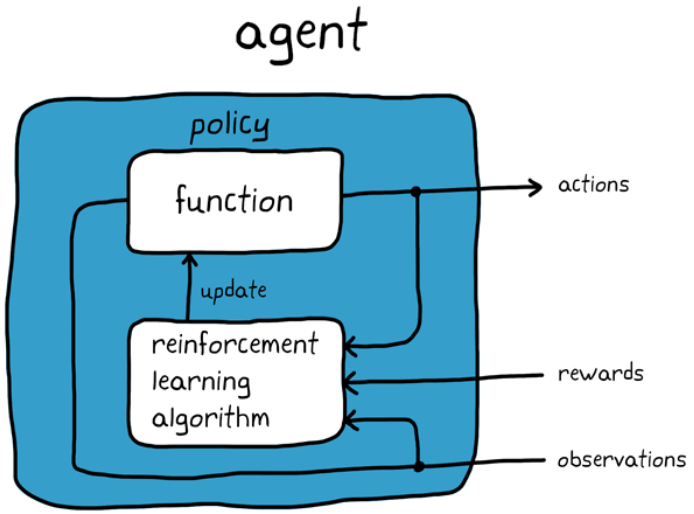

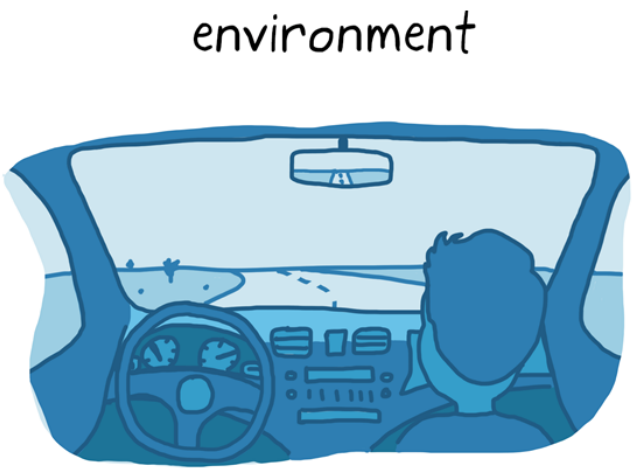

Source: Reinforcement learning – Wikipedia

Source: What Is Reinforcement Learning? - MATLAB & Simulink

The ML Lifecycle

How does it work?

How does it work?

Collect

- Unknown sender
- Spelling mistakes
- Lowercase names
- Promise of free winning
- Sense of urgency
- Strange links
- Request for payment
- The label: Malicious or Not Malicious

Prepare

- Transform and join data tables
- Handle errors, outliers and missing values
- Convert strings to numbers
- Further data transformations

Train

- Look at similar emails – are they malicious or not?
- Determine a set of rules: if this and that, then it is likely malicious. Otherwise…
- Calculate a risk score as a weighted sum of the inputs

Evaluate

- What percentage did you guess right?
- How many malicious emails did you not notice?
- How many emails did you mistakenly believe to be malicious?

The Machine Learning Lifecycle

Source: Cross-industry standard process for data mining - Wikipedia

ML Algorithms

Machine Learning Algorithms

Supervised learning / Classification

- Logistic Regression
- Decision Tree
- k-nearest Neighbors (kNN)
- Support Vector Machines (SVM)
- Ensembles (Random Forest, Gradient Boosting)

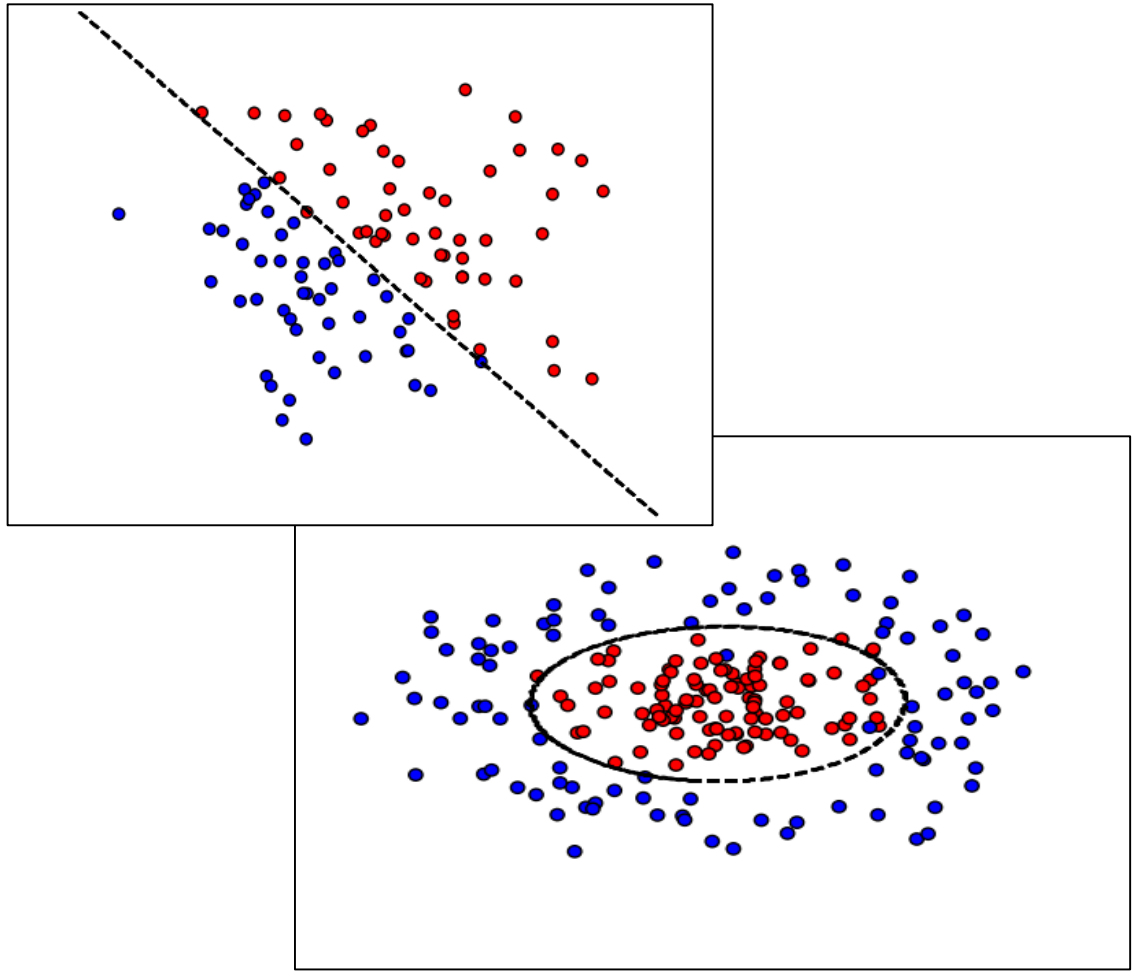

Machine Learning Algorithms

Supervised learning / Regression

- Linear Regression
- Decision Tree
- k-nearest Neighbors (kNN)
- Support Vector Machines (SVM)
- Ensembles (Random Forest, Gradient Boosting)

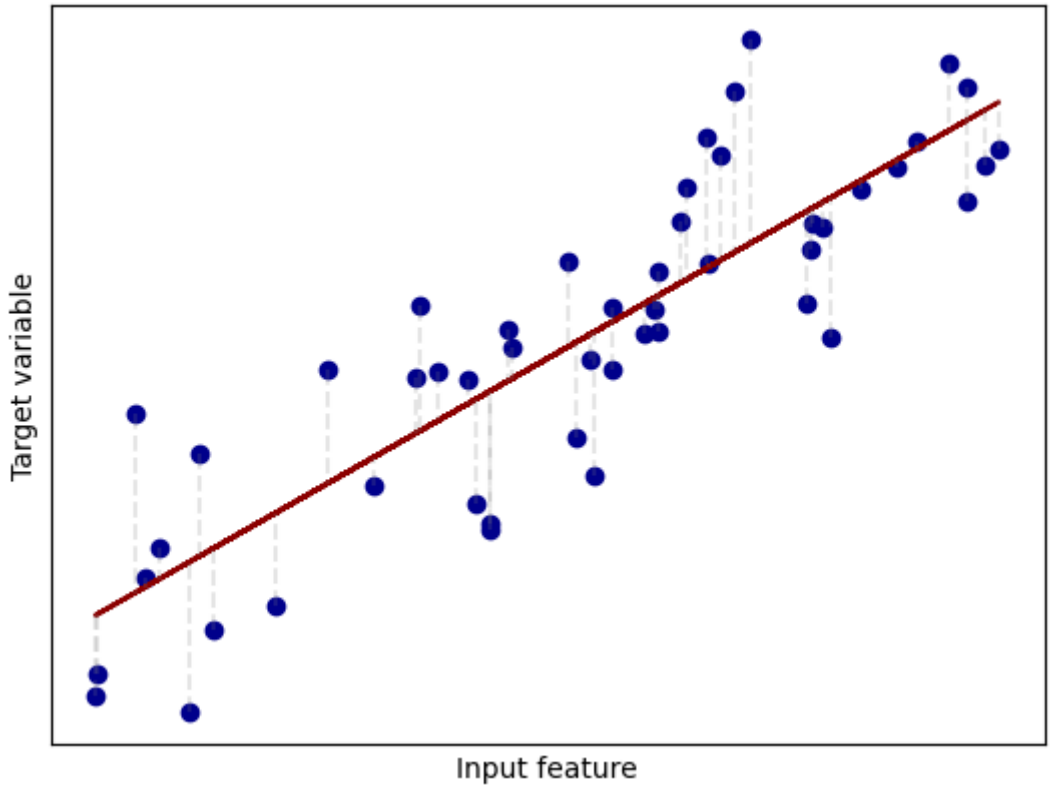

Machine Learning Algorithms

Unsupervised learning / Clustering

- K-means
- Hierarchical
- Density based (e.g. DBSCAN)

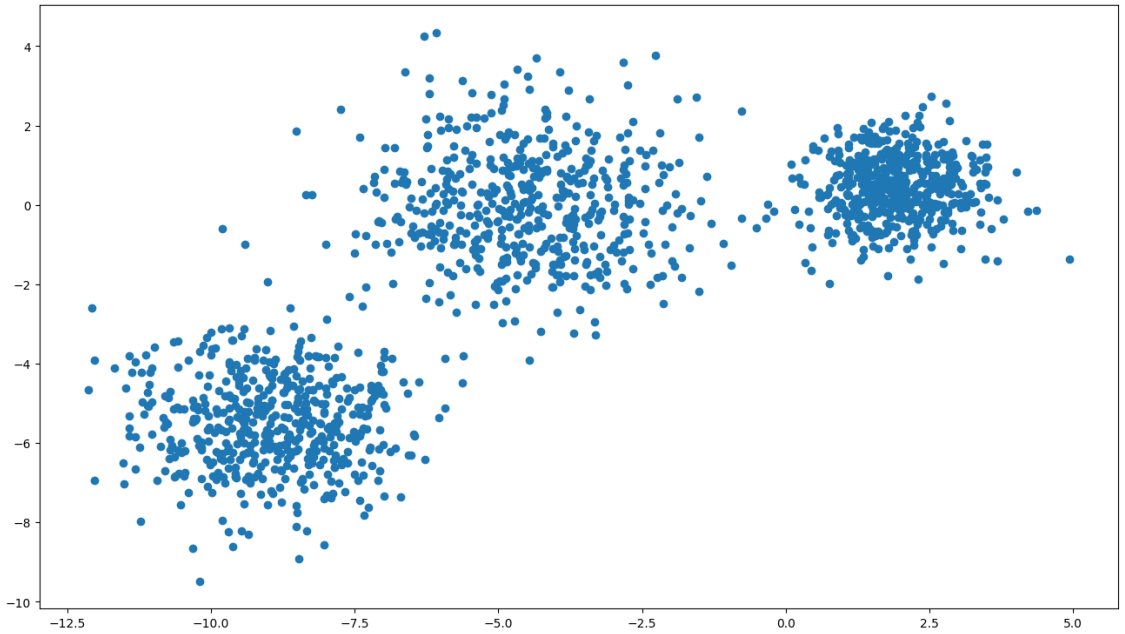

Machine Learning Algorithms

Unsupervised learning / Anomaly Detection

- Statistical Outlier Detection
- Isolation Forest
- One-class SVM
- Autoencoder

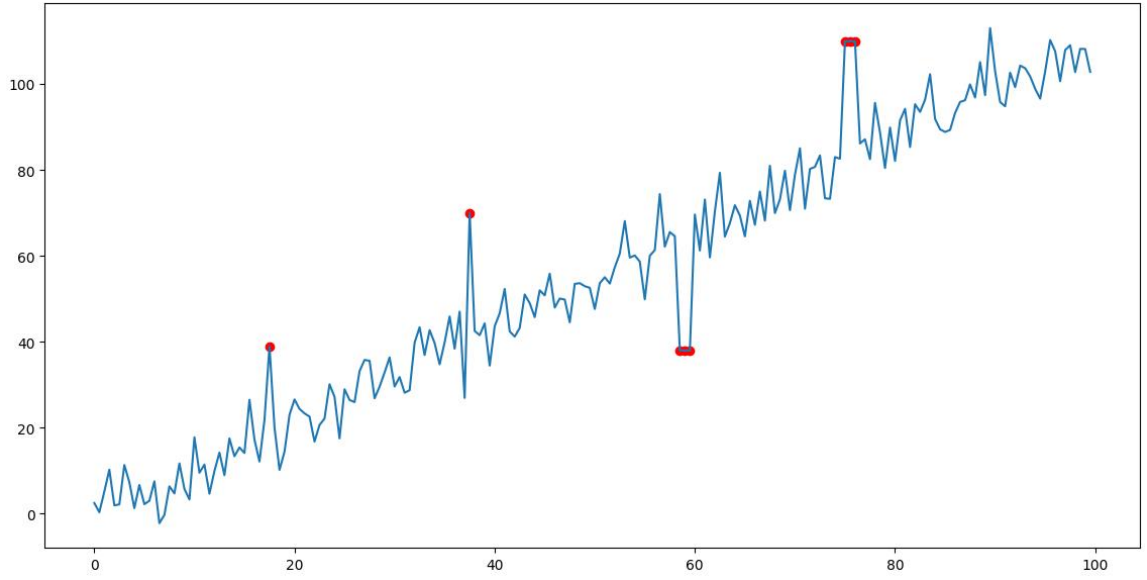

Machine Learning Algorithms

Unsupervised learning / Dimensionality Reduction

- PCA
- Factor Analysis
- Autoencoder

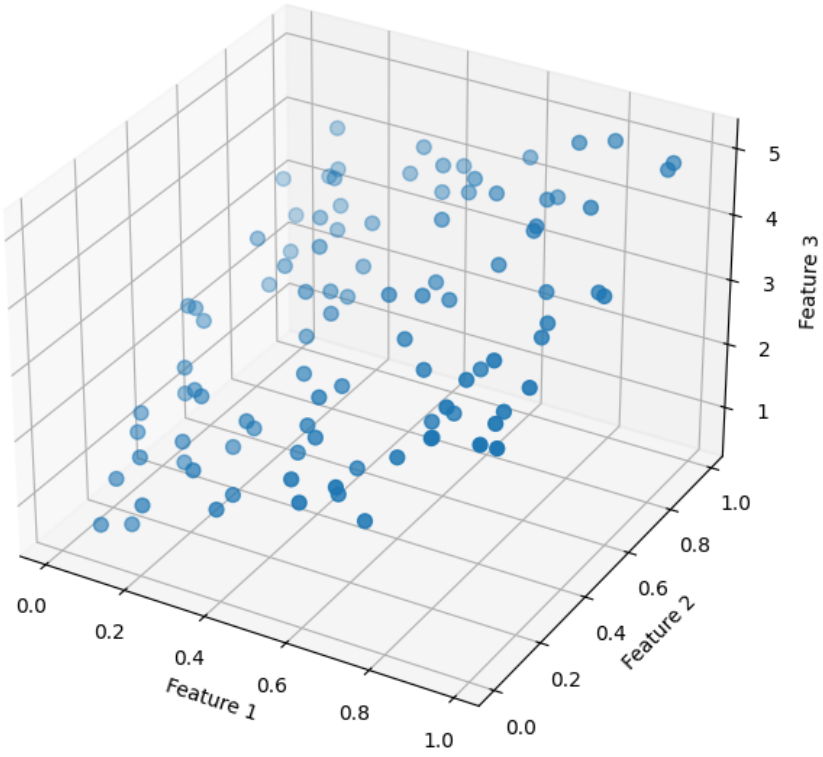

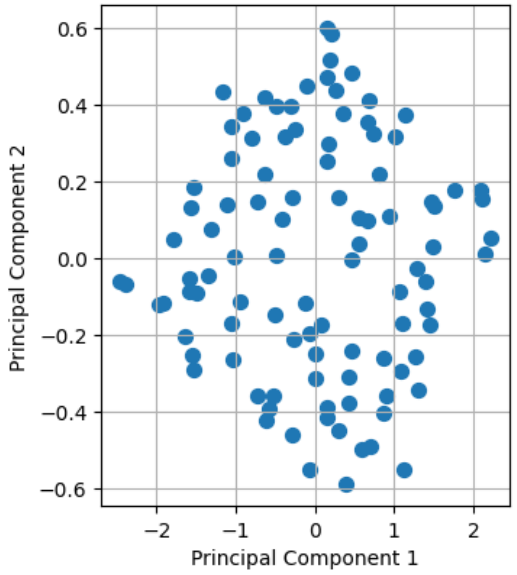

Machine Learning Algorithms

Unsupervised learning / Recommender Systems

- Collaborative filtering
- Content-based filtering
- Matrix factorization

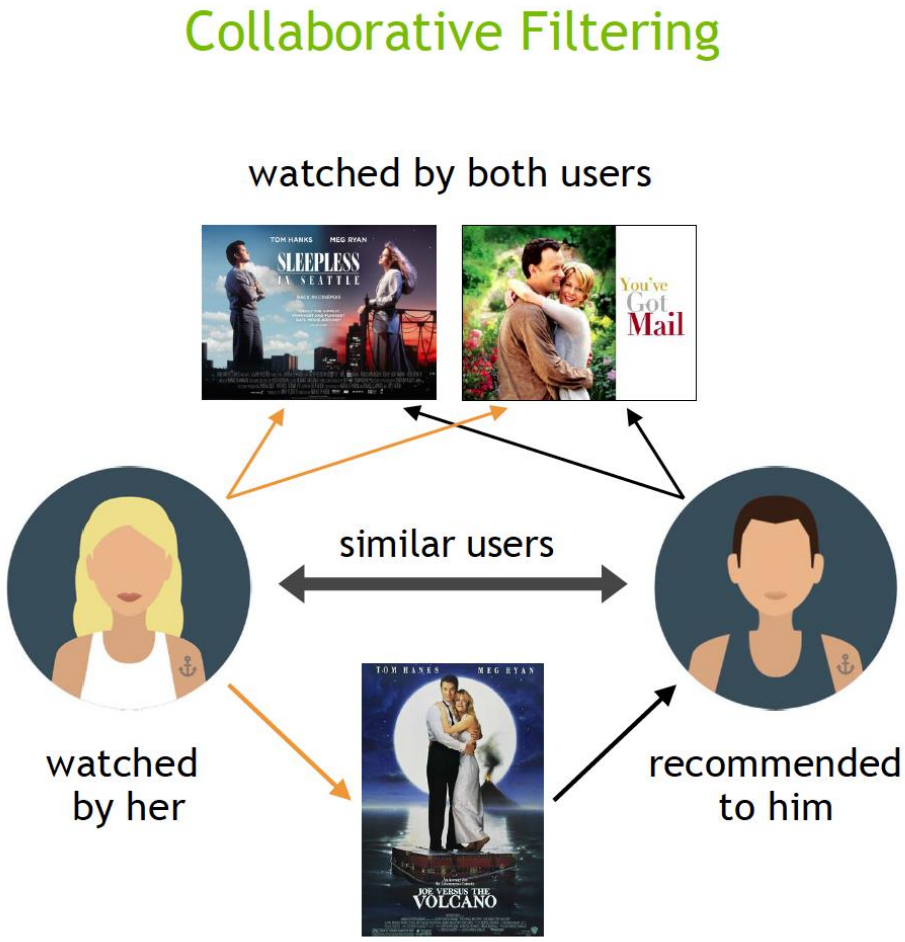

Source: What is a Recommendation System? | Data Science | NVIDIA Glossary

Model Training

Logistic Regression

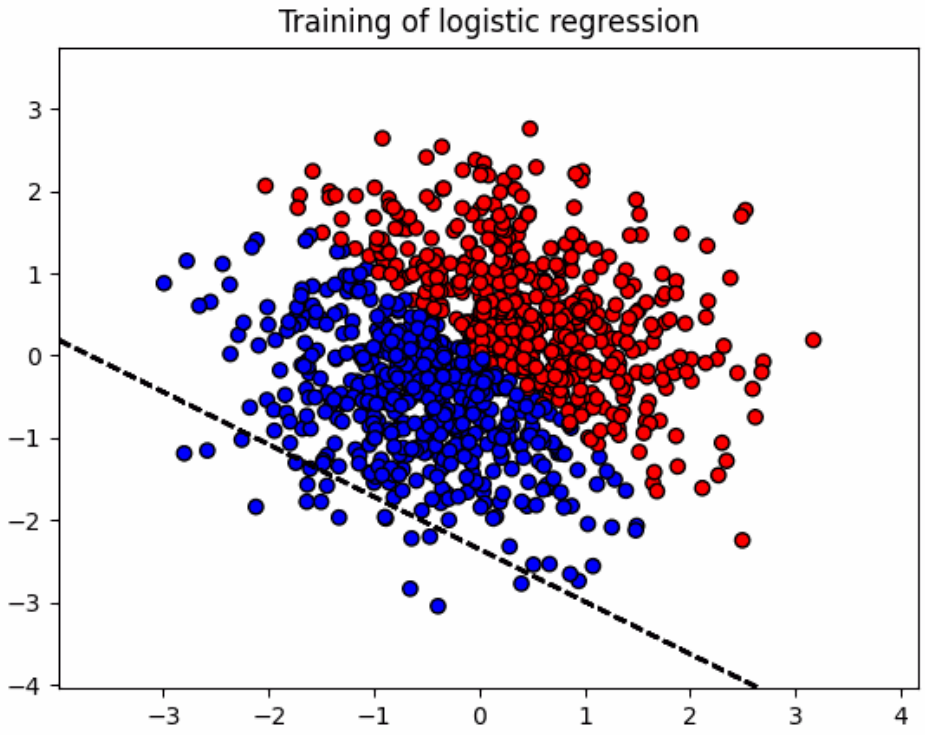

K-means Clustering

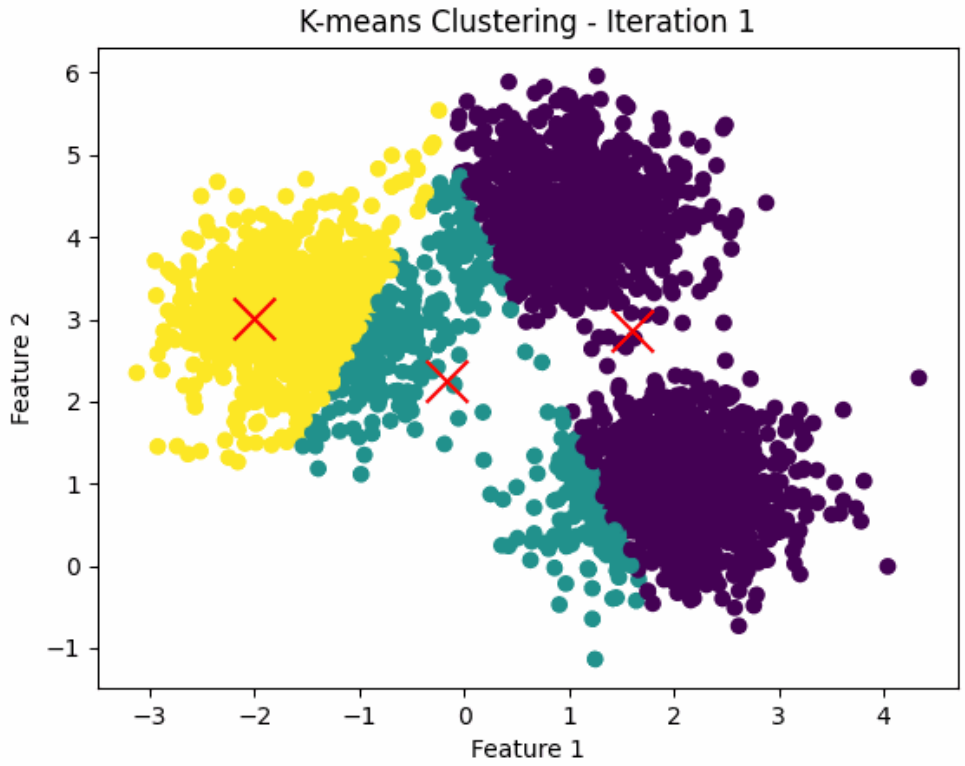

The Goal

Generalization

The goal of ML is to develop models that perform well on unseen data.

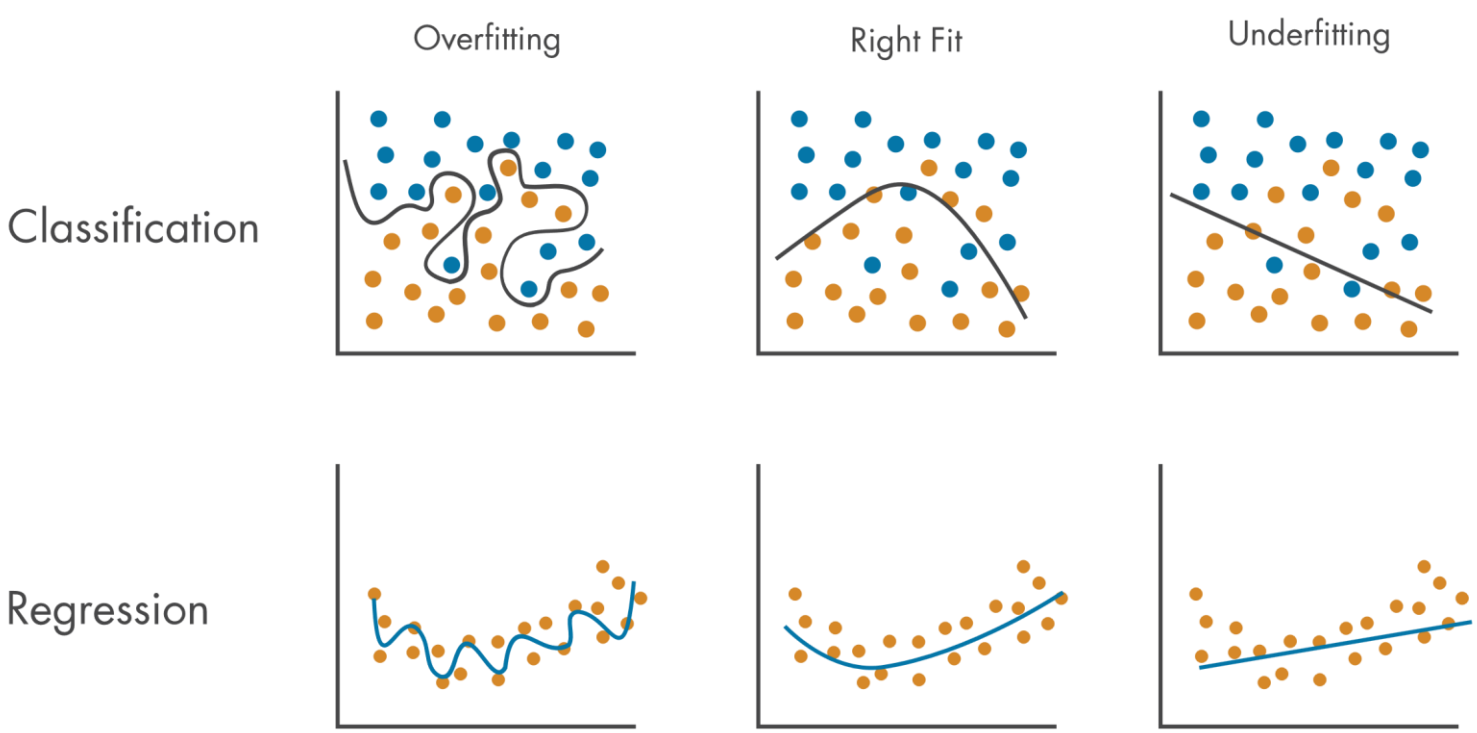

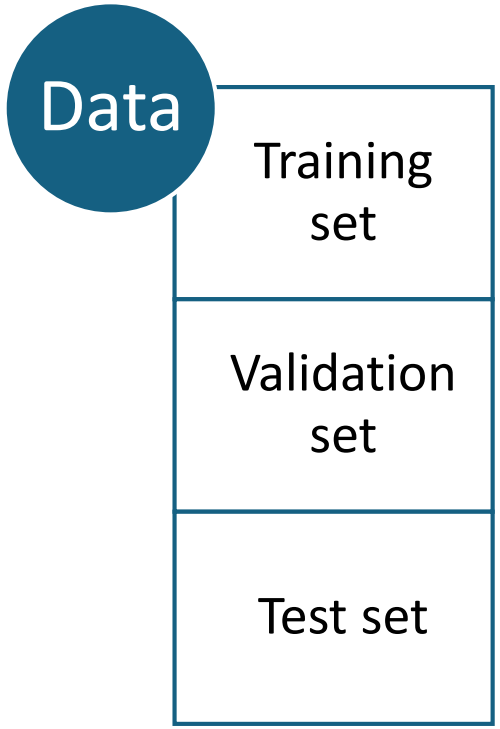

Source: Overfitting - MATLAB & Simulink

Deployment

More than 80% of data science projects never make it to production

Pitfalls preventing deployment

- Lack of engagement
- No real business value
- Data issues
  - Availability
  - Quality
  - Regulation
- Implementation issues
  - Technical Integration
  - Lack of expertise
  - Budget

Pitfalls after deployment

- Bad model
  - Bias
  - Poor performance
- Different data/environment
- Bad model-based decisions
- External factors

To do

- Consider deployment from the beginning of a project
- Talk to the business
- Model monitoring
- Be ready to intervene

Deep Learning

What is Deep Learning?

A subset of machine learning inspired by the human brain, utilizing deep neural networks to learn from data. A neural network is basically a sequence of nonlinear mathematical functions.

Model architecture outlines the sequence and connectivity of layers and neurons, along with the functions they perform.

Elements of a Neural Network:
- Layers
- Neurons (Nodes)
- Weights and biases
- Activation functions

Model Training:
- Gradient Descent
- Backpropagation

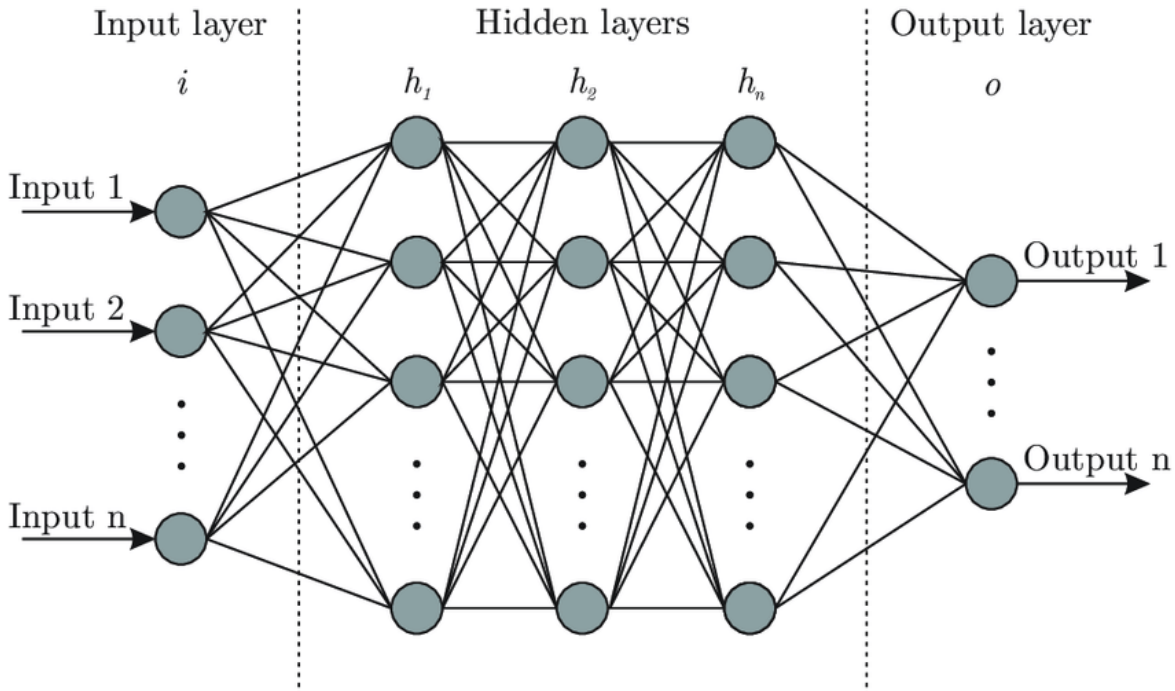

Source: Designing Your Neural Networks. A Step by Step Walkthrough | by Lavanya Shukla | Towards Data Science

What is Deep Learning?

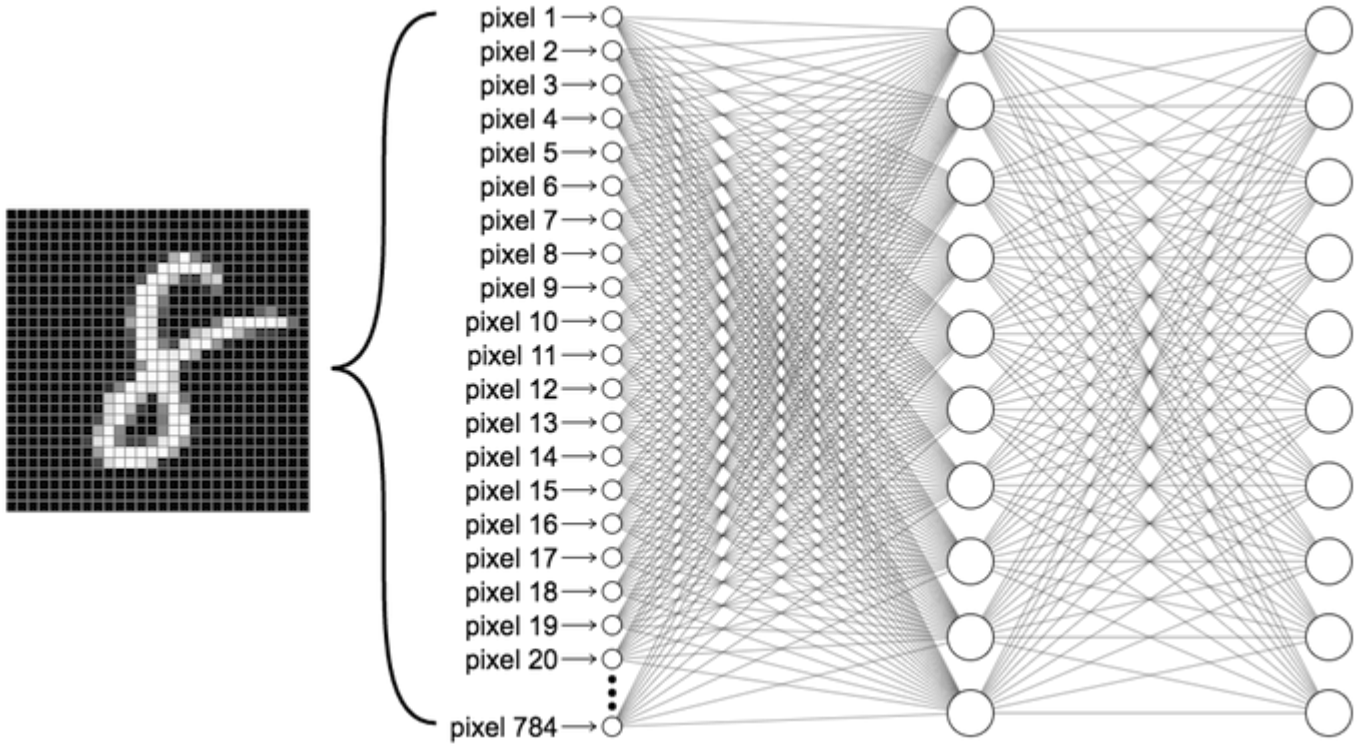

Handwritten Digit Recognition with a Feedforward Neural Network

Source: Designing Your Neural Networks. A Step by Step Walkthrough | by Lavanya Shukla | Towards Data Science

What is Deep Learning?

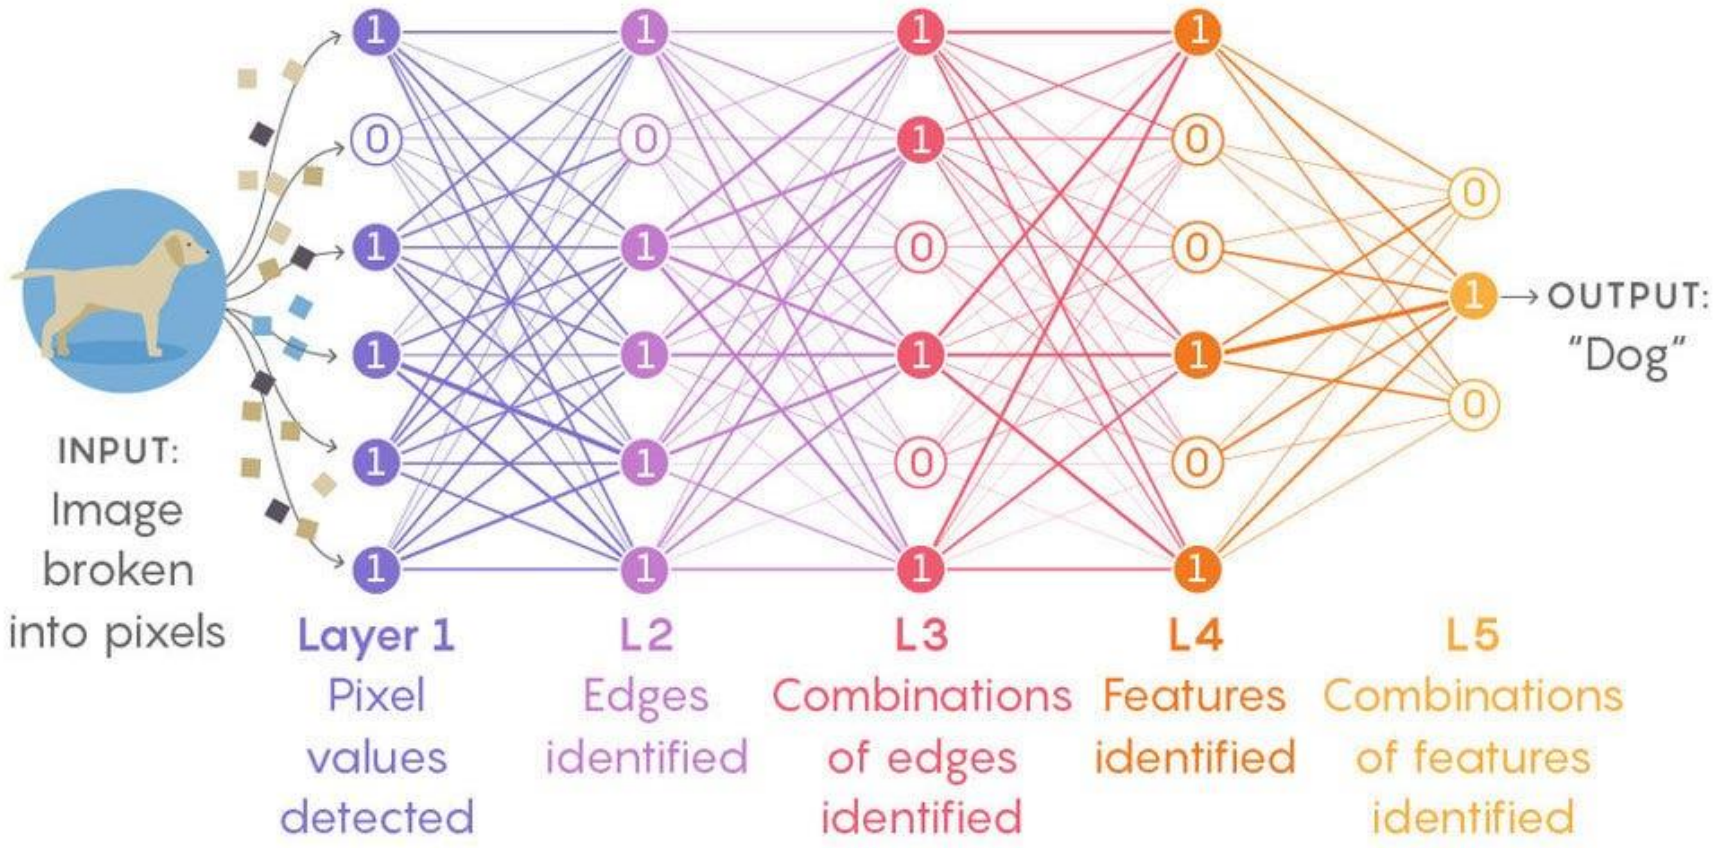

Source: Designing Your Neural Networks. A Step by Step Walkthrough | by Lavanya Shukla | Towards Data Science

A Brief History of AI with Deep Learning

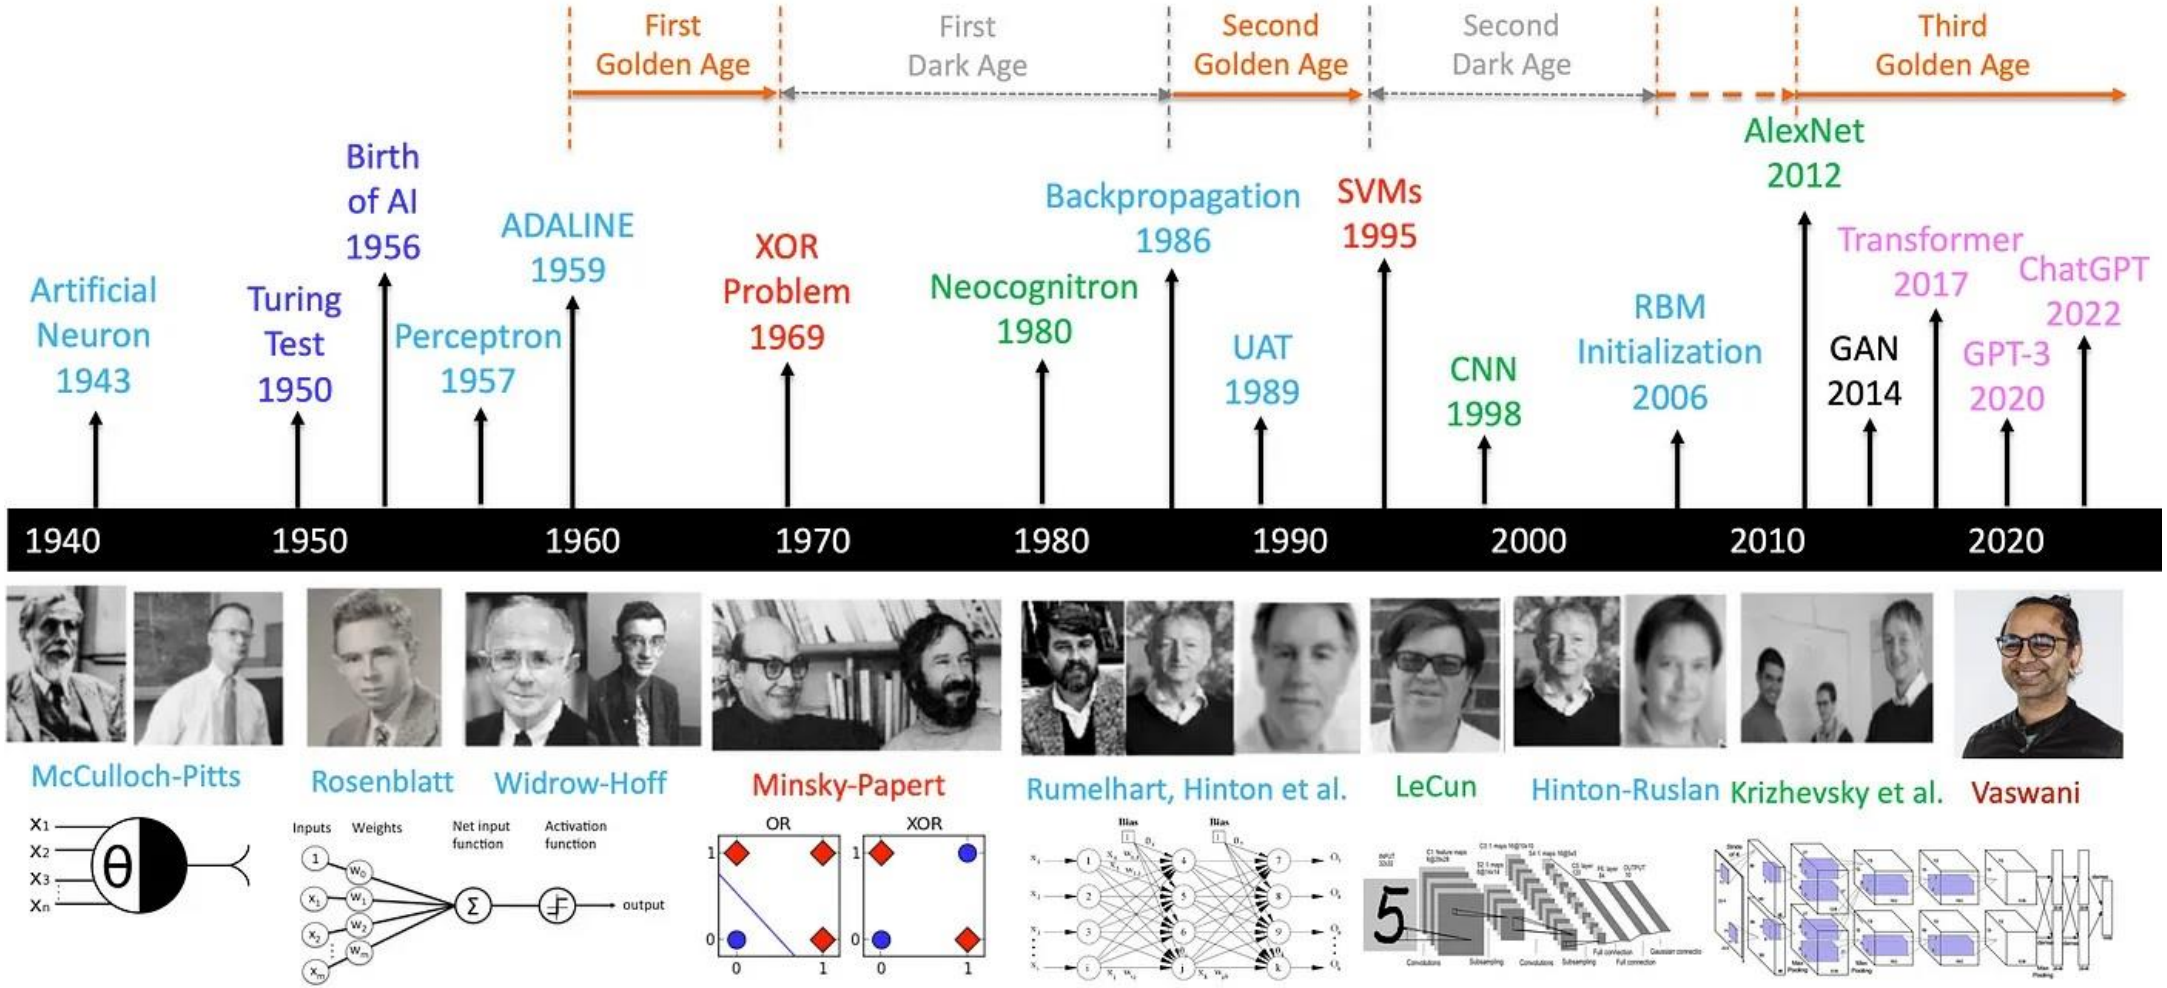

Golden Age from 2012 – Data, Compute, Algorithms

Source: A Brief History of AI with Deep Learning | by LM Po | Medium

Sequence Models

Time Series	Text	Audio	Architectures
• Forecasting • Anomaly Detection • Classification • Imputation	• Text classification • Named Entity Recognition • Sentiment Analysis • Text Summarization • Machine Translation • Text Generation	• Text-to-speech • Speech-to-text • Music Generation	• RNN • LSTM • GRU • Encoder-decoder • Transformer

Deep Learning for Text

One-hot representation

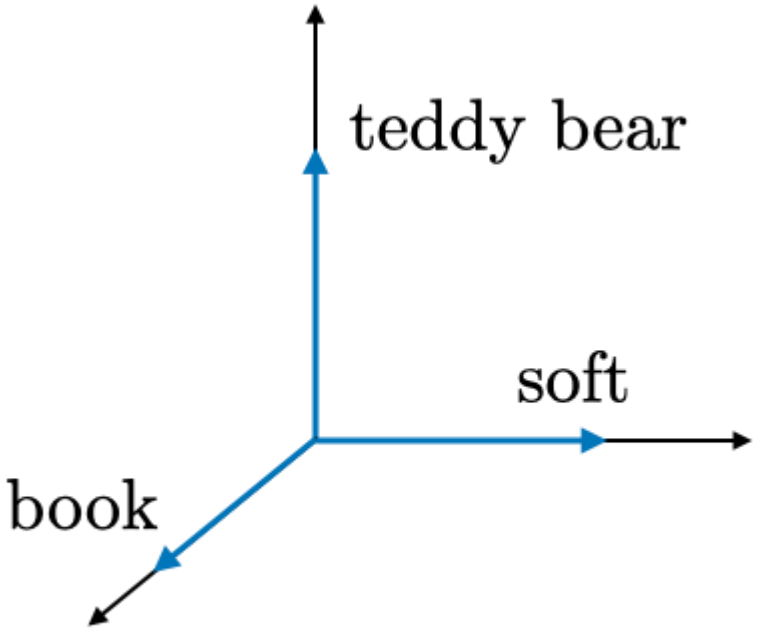

Vector Embedding

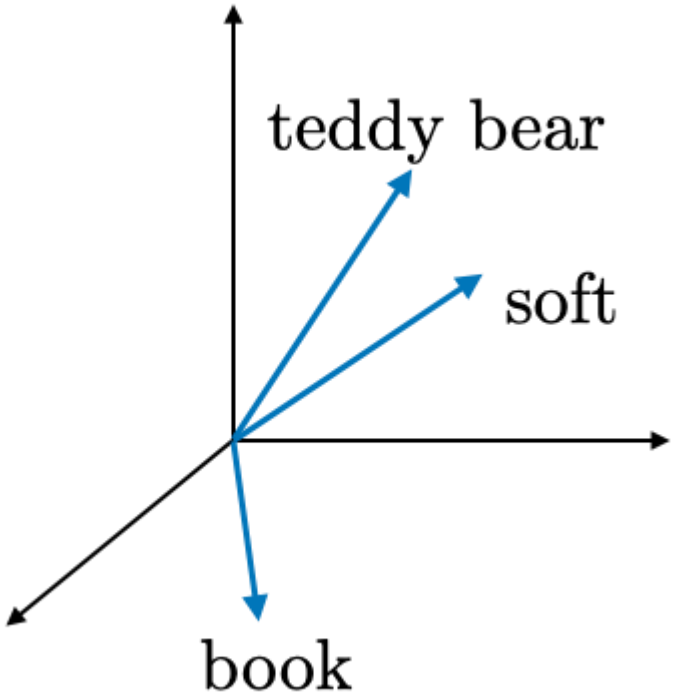

Cosine Similarity

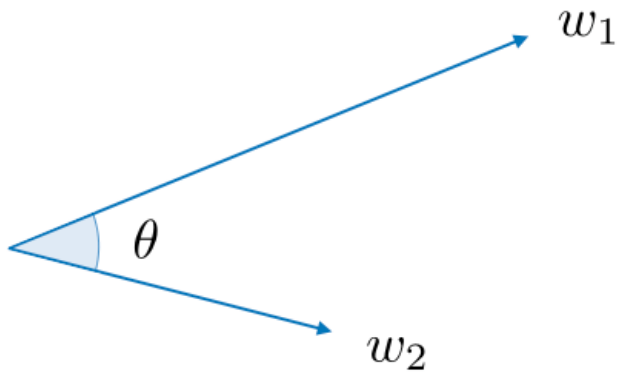

$$similarity = \frac{\omega_1 \omega_2}{\|w_1\|\|w_2\|} = \cos(\theta)$$

Embedding: words (tokens) are represented as multidimensional vectors. Directions and distances describe the relationships between words. Embeddings are optimized to best model these relationships.

Source: CS 230 - Recurrent Neural Networks Cheatsheet

Deep Learning for Vision

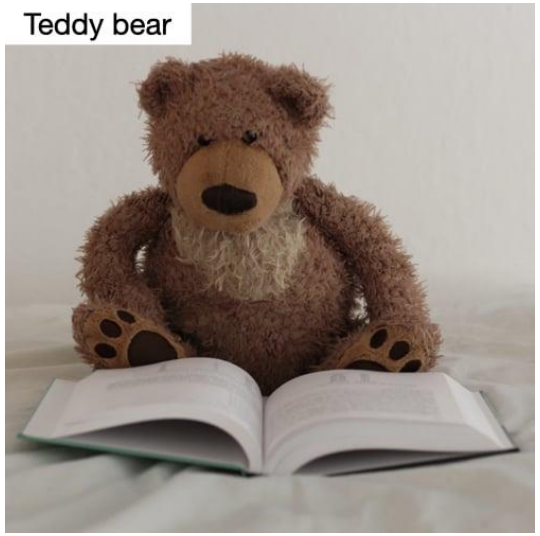

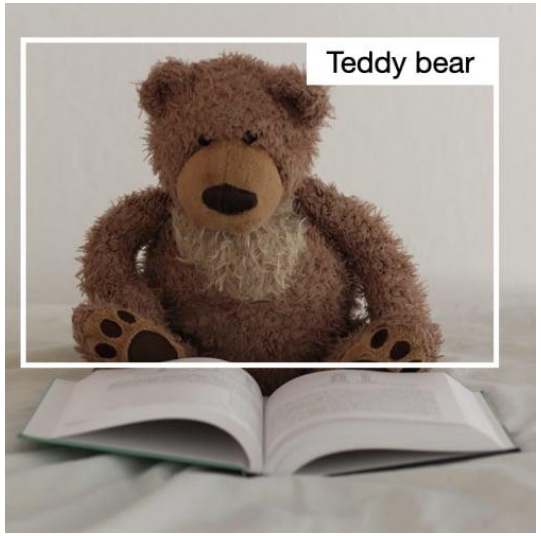

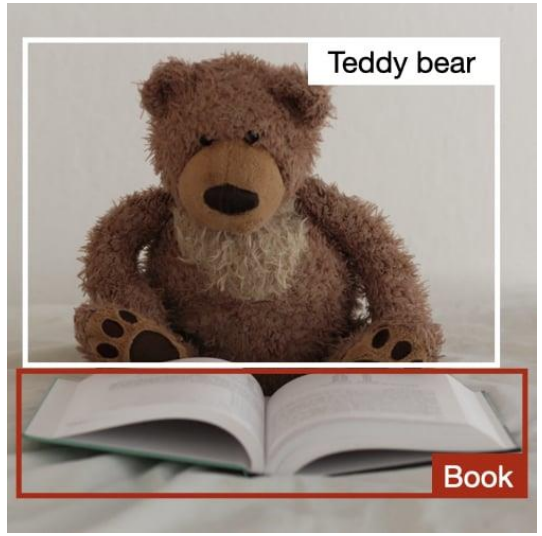

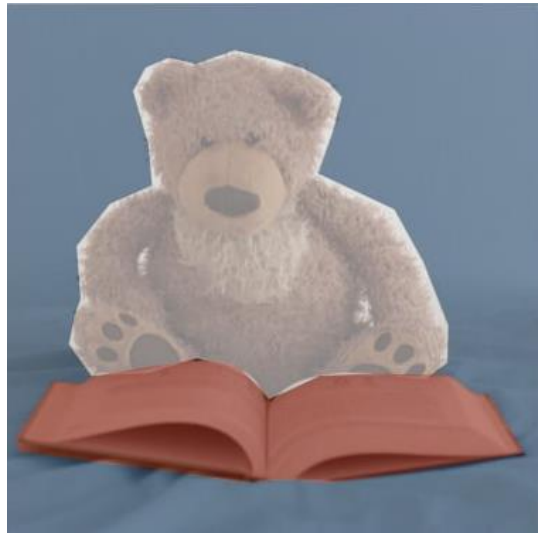

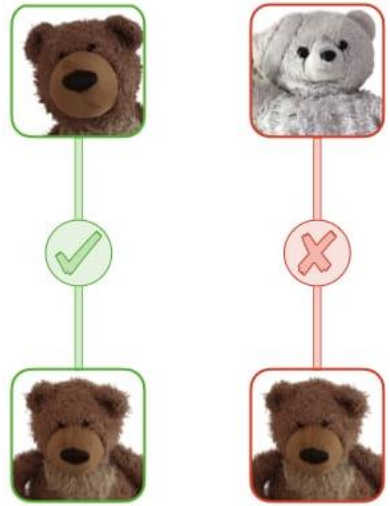

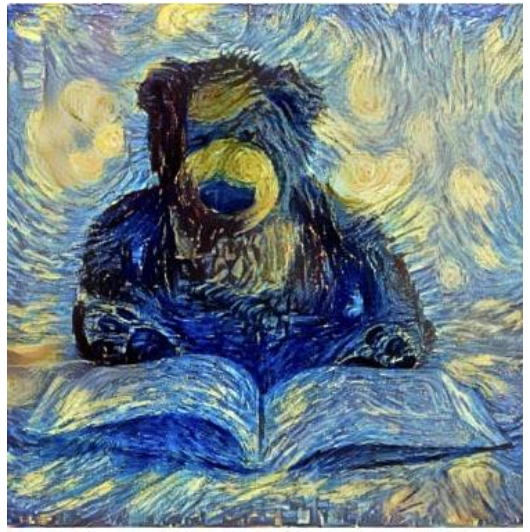

Source: CS 230 - Convolutional Neural Networks Cheatsheet

Computer Vision

Image

- Image recognition
- Object detection
- Semantic /Instance Segmentation
- Optical Character Recognition
- Facial Recognition

Generative

- Image Enhancement
- Style Transfer
- Image Captioning
- Video Manipulation
- Image/Video Generation

Video / 3D

- Object Tracking
- Pose estimation
- Depth Estimation
- 3D Reconstruction
- SLAM

Architectures

- Traditional CV
- Deep Learning:
  - CNN
  - Autoencoder
  - GAN
  - Diffusion Model
  - Vision Transformer
  - Multimodality

What are LLMs?

- **Deep learning models** designed and optimized specifically for conversations.
- **Transformer:** groundbreaking model architecture based on the **attention** mechanism – the word representation is influenced by the context
- **Pre-training** on large data, **post-training** on high quality data with methods like RLHF
- **System prompt:** instructions to determine how the AI model responds to the **user prompt**
- The trained weights determine the model's behaviour by storing the model's understanding of the world (memory). But this is not human intelligence – more like "autocorrect on steroids"

- Learn more:
  - 3Blue1Brown – Neural networks
  - Andrej Karpathy – Deep Dive into LLMs like ChatGPT
  - Stanford CS229 I Machine Learning I Building Large Language Models (LLMs)

Deep Learning versus Machine Learning

Deep Learning:

- Built-in feature extraction
- Better architectures for text and vision:
  - Word embedding for text
  - Convolution for images
  - Sequence modelling for time series and text
- Performs well on large amounts of data and complex problems
- Computationally expensive, but great parallelism (GPUs)

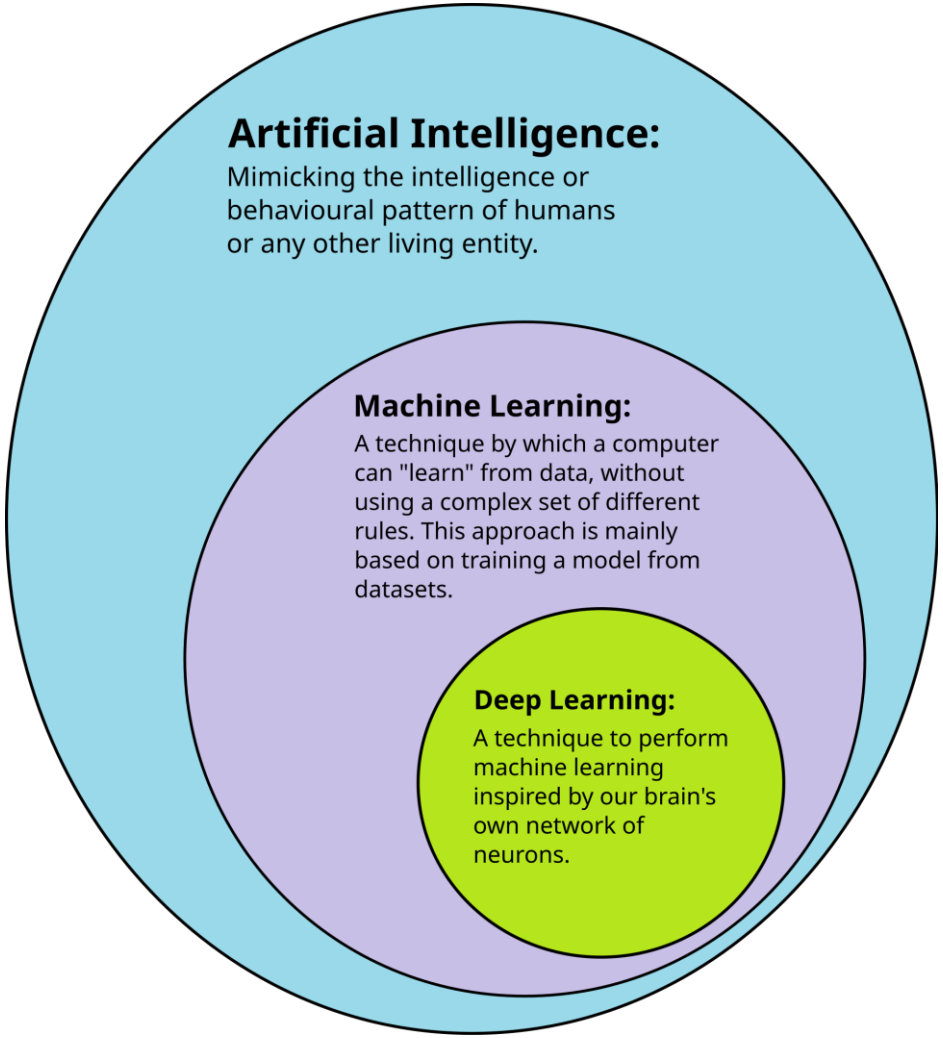

Source: Deep learning - Wikipedia

Mathematical Optimization

Maximize/minimize a target function (cost, time), or find a solution if it exists, with certain constraints, limited choices or limited computing resources.

- Search Algorithms
  - State space search
    - Shortest path
    - Minimum Spanning Tree
  - Local search
    - Hill climbing
    - Simulated Annealing
    - Travelling Salesman

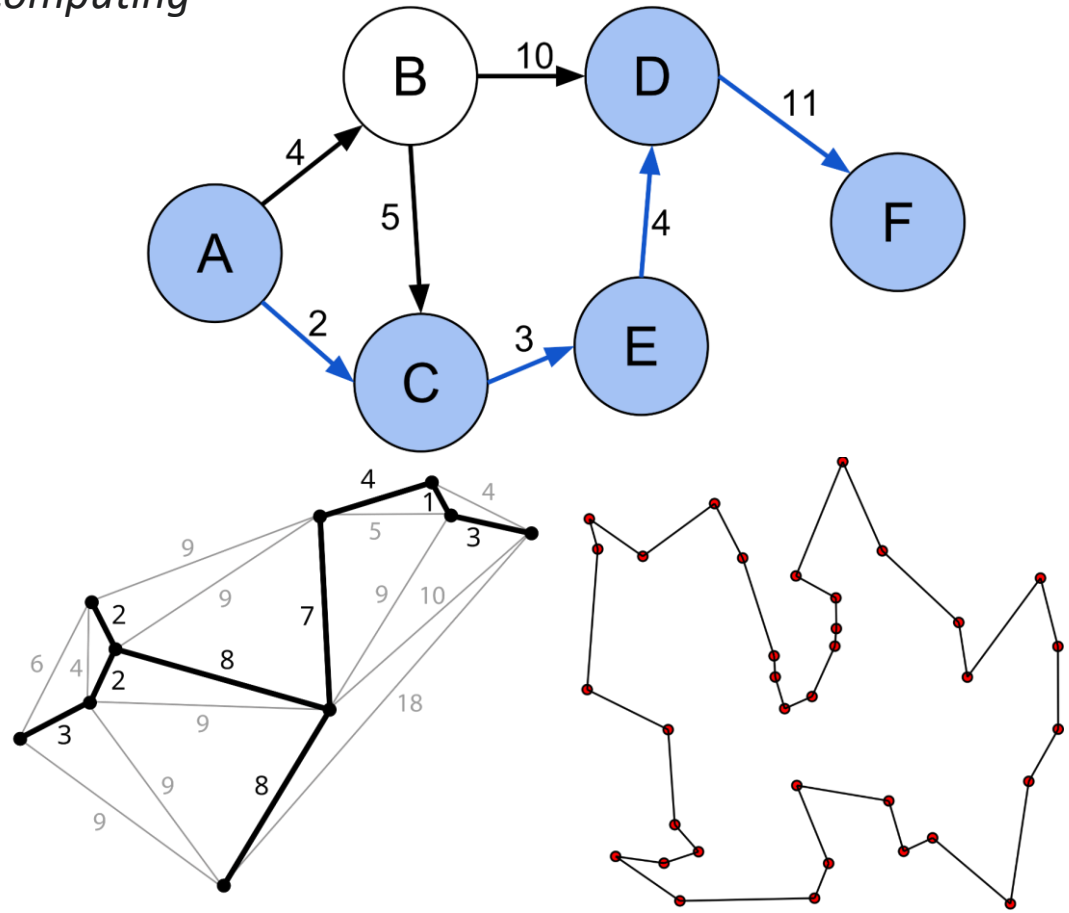

Source: Wikipedia (Shortest path problem, Minimum spanning tree, Travelling salesman problem)

Mathematical Optimization

- Evolutionary Algorithms
- Swarm Intelligence
- Logic (Knowledge Representation)
- Constrained Optimization
  - LP, NLP, IP, QP
  - Constraint Satisfaction
  - Knapsack Problem
- Stochastic Optimization
  - Monte Carlo Simulation

Learn more: CS50's Introduction to Artificial Intelligence with Python

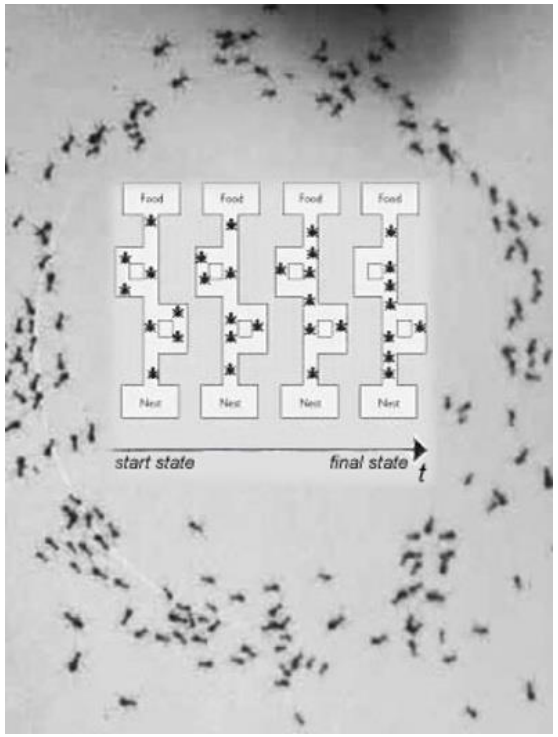

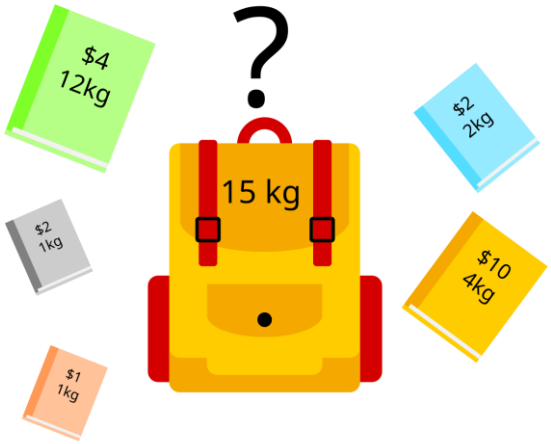

$$\max c^T x$$
$$subject\ to\ Ax \leq b$$
$$and\ x \geq 0$$

Source: Wikipedia (Ant colony optimization algorithms, Knapsack problem)

Responsible AI

- Data Privacy & Transparency
- Fairness & Inclusivity
- Reliability & Safety
- Explainability
- Accountability
- Sustainability

Tools

Skills for Data Science

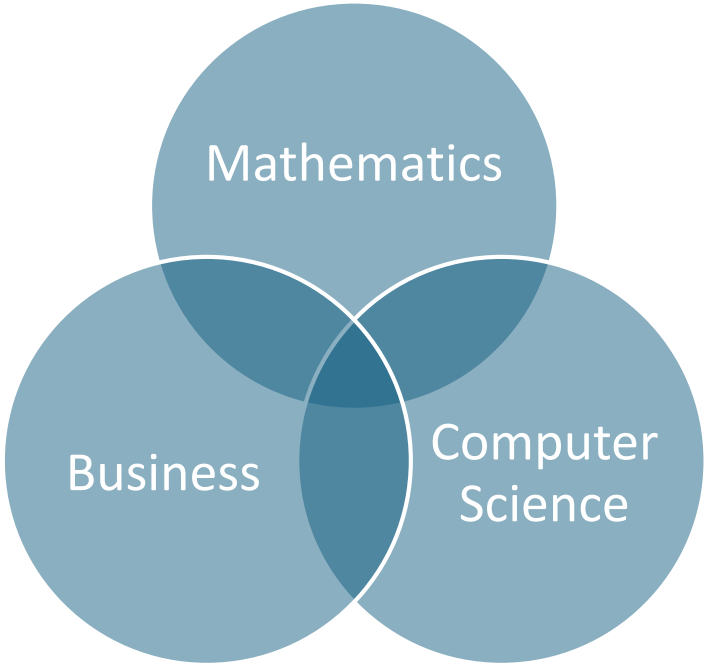

- **Computer Science**
  - Python / R
  - SQL
  - DevOps
  - Cloud
- **Business**
  - Communication
  - Business Acumen
  - Storytelling
- **Mathematics**
  - Probability Theory
  - Statistics
  - Linear Algebra
  - Calculus
- **Sherlock Personality**
  - Observation
  - Deduction
  - Knowledge
  - Curiosity

Python for Data Science

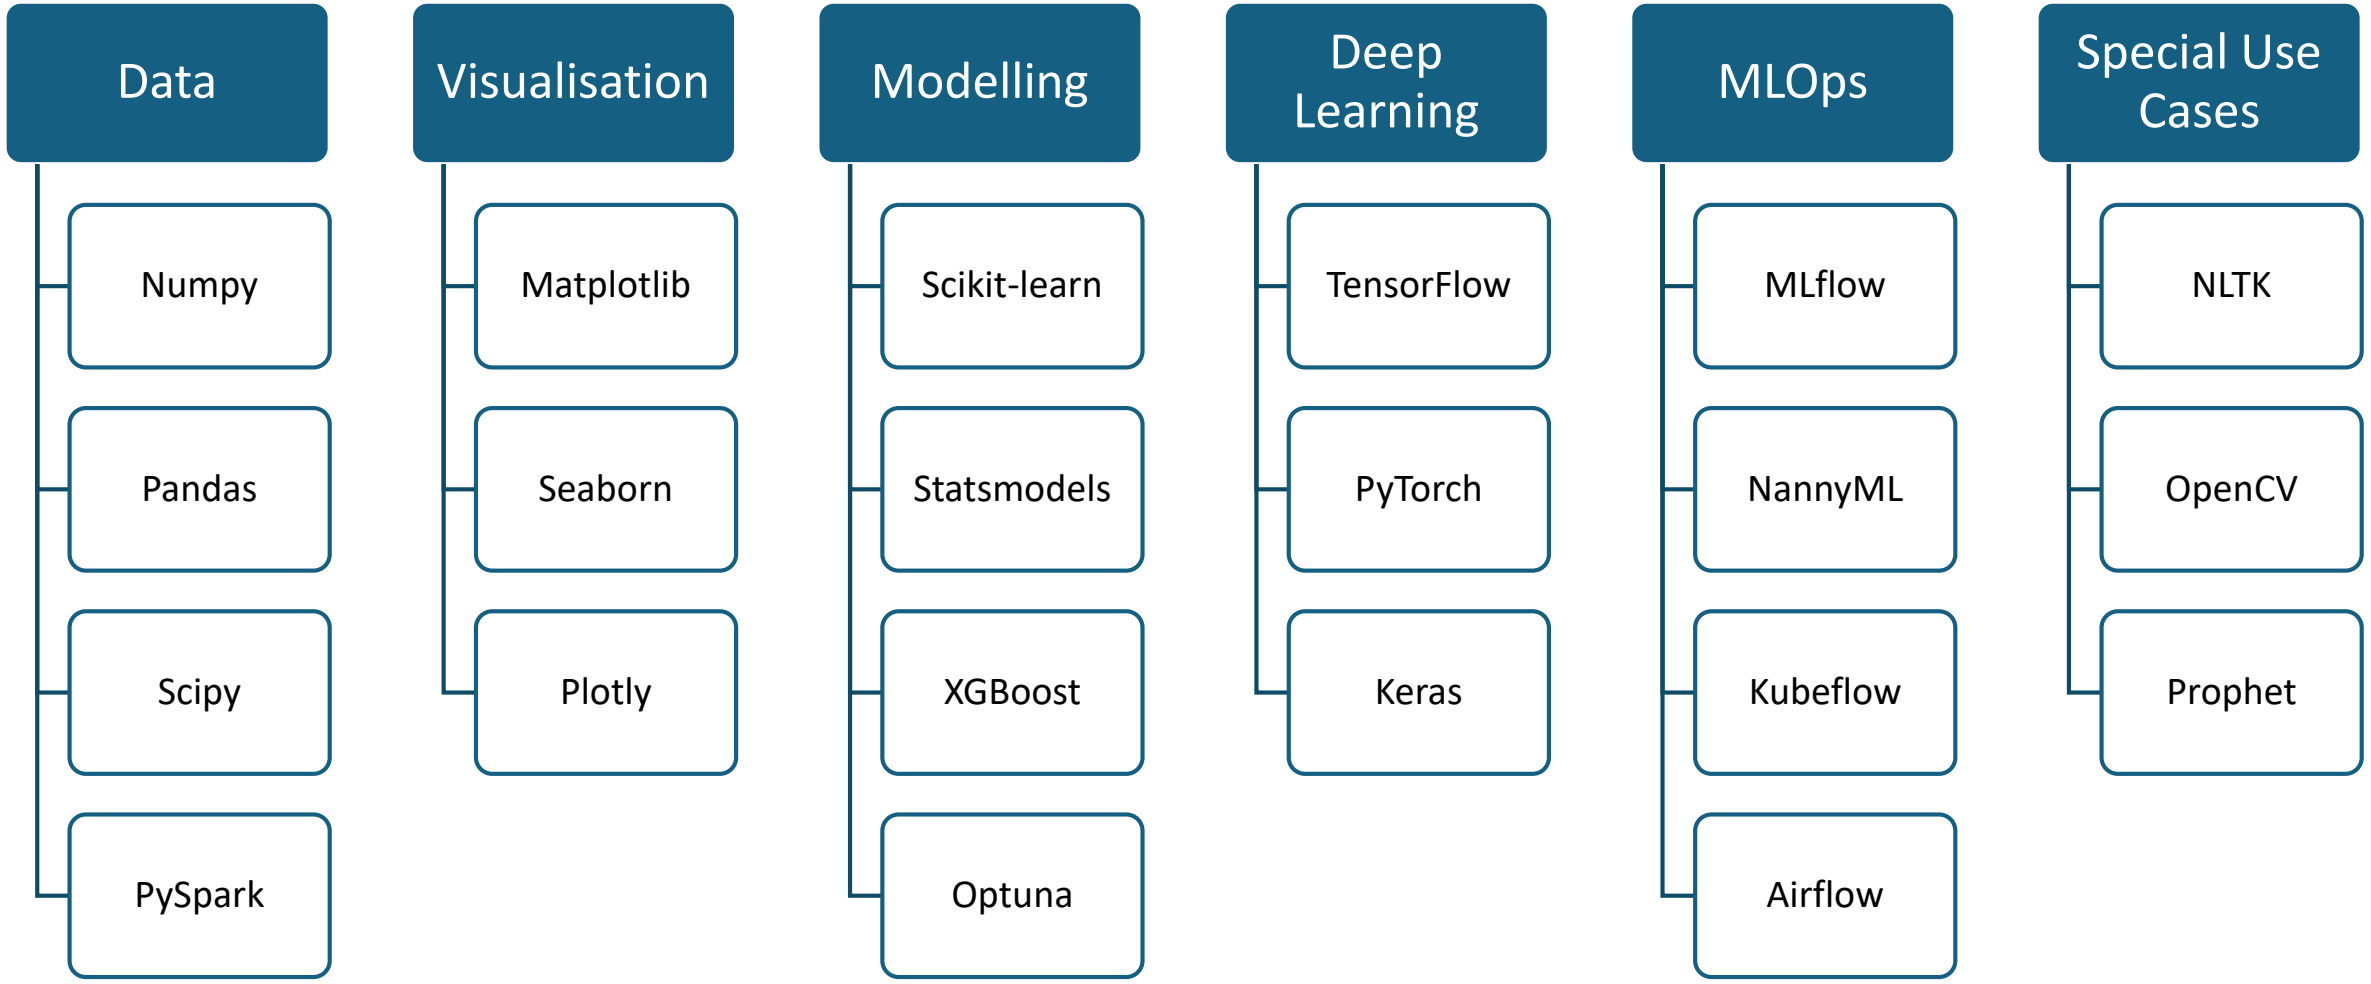

Cloud Platforms for Data Science

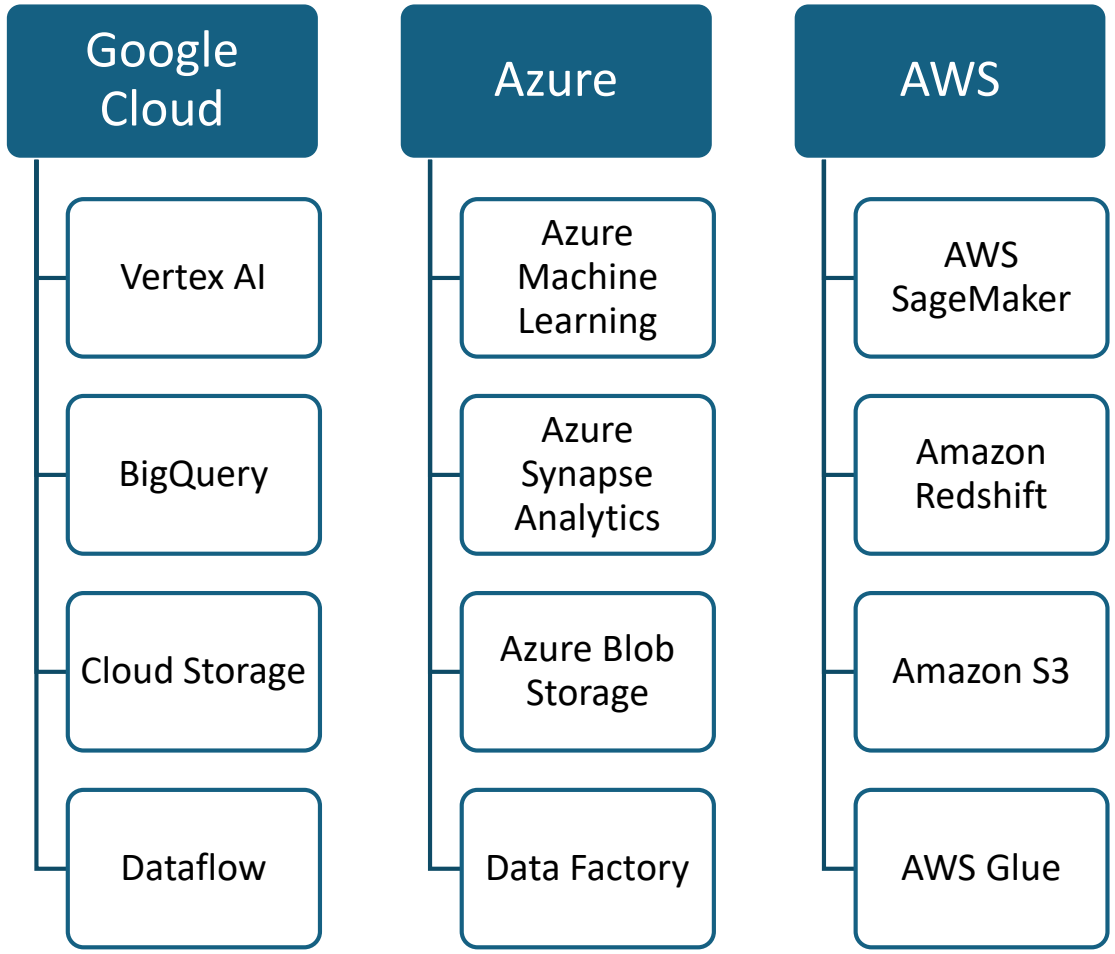

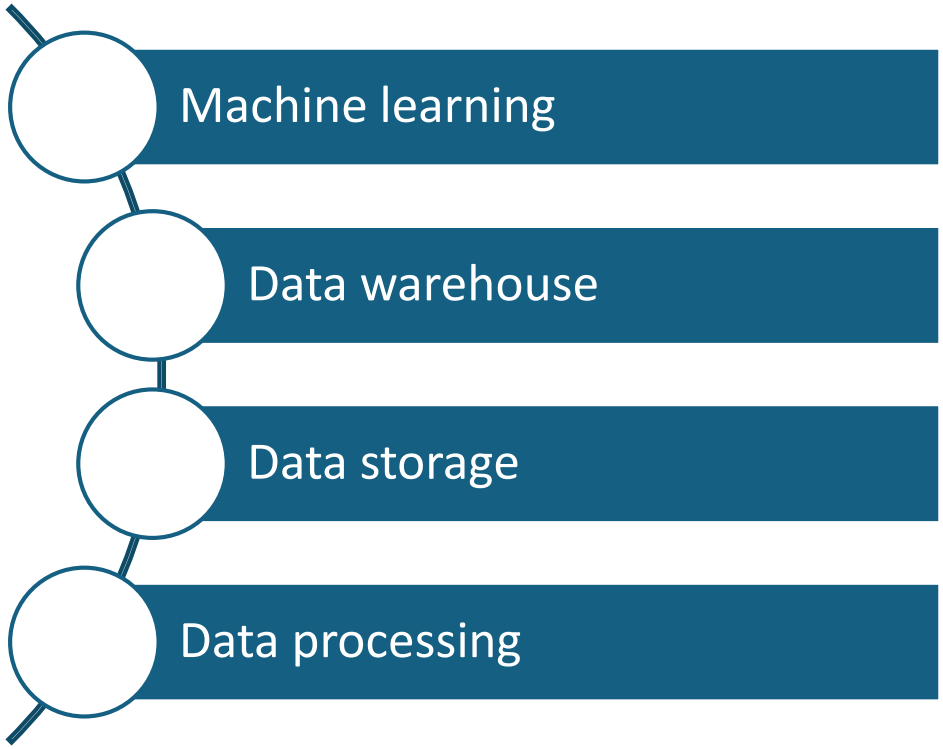

Introduction to Data Science

Thank you for your attention!

Your feedback would be much appreciated:

Any Questions?

Gergely Zsombor Haász

haasz.zsombi@gmail.com