

A Practical Introduction to Data Science

Part 4

Advanced Supervised Learning



Gergely Zsombor Haász



haasz.zsombi@gmail.com

Course Agenda

- I. Introduction to Data Science
- II. Business and Data Understanding
- III. Introduction to Supervised Learning
- IV. Advanced Supervised Learning
- V. Unsupervised Learning
- VI. Time Series Analysis
- VII. Deep Learning
- VIII. Machine Learning Operations

Advanced Supervised Learning

Tree Ensembles

Overfitting and Generalization

Hyperparameter Optimization

Tree Ensembles

Random Forest

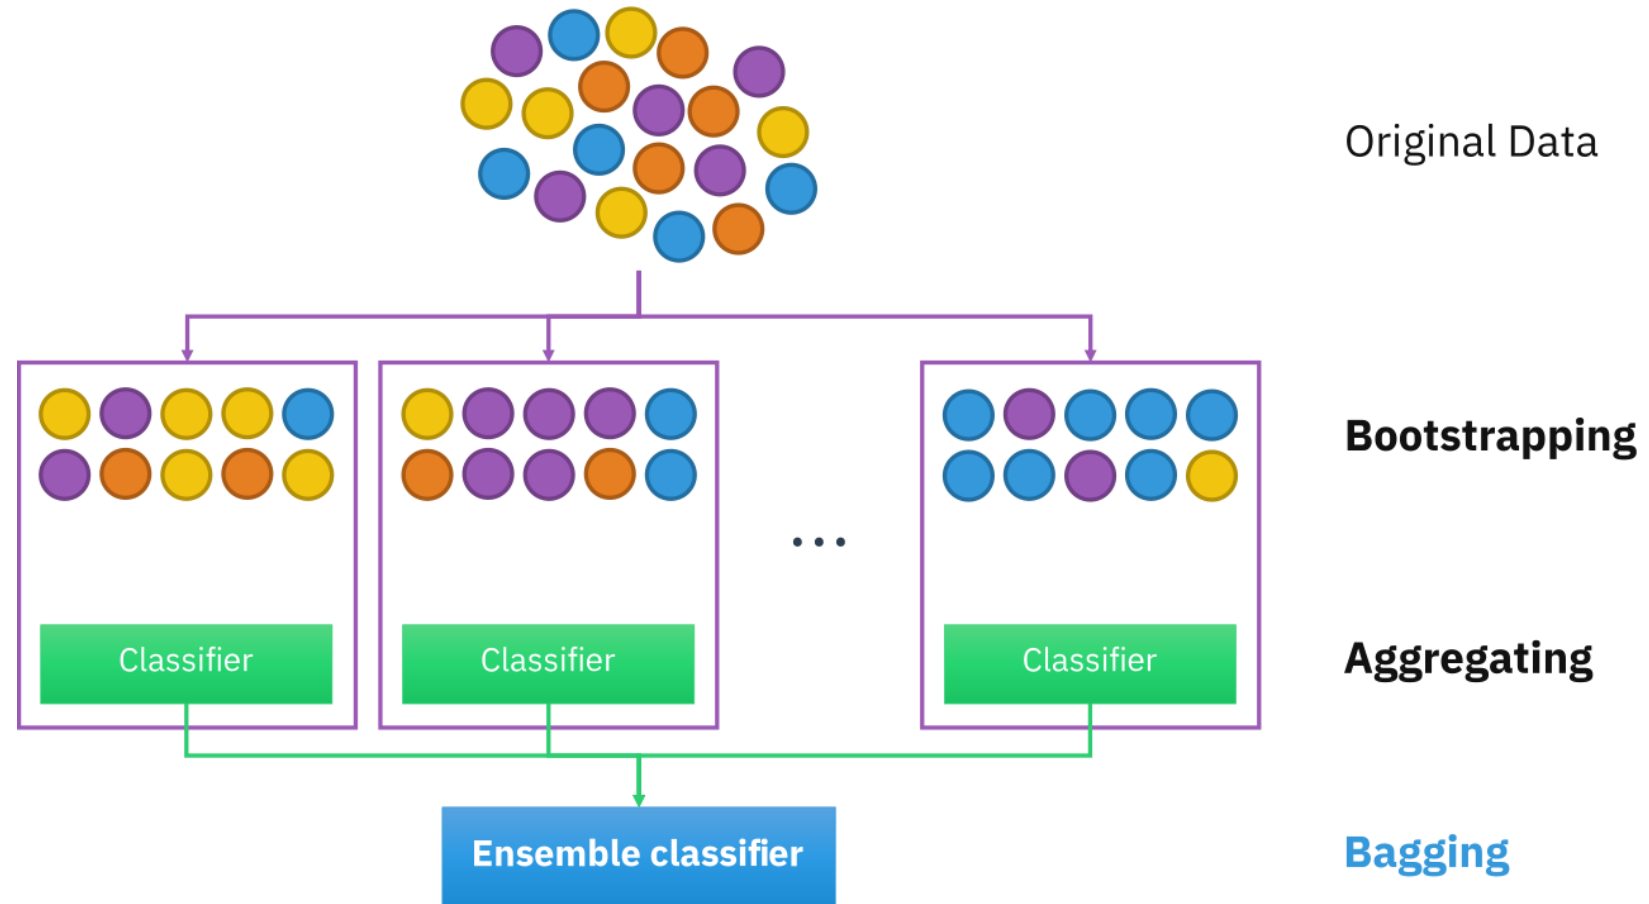
- Base learner: Decision Tree
- Bagging + random subspace method

Advantages:

- More accurate
- More robust

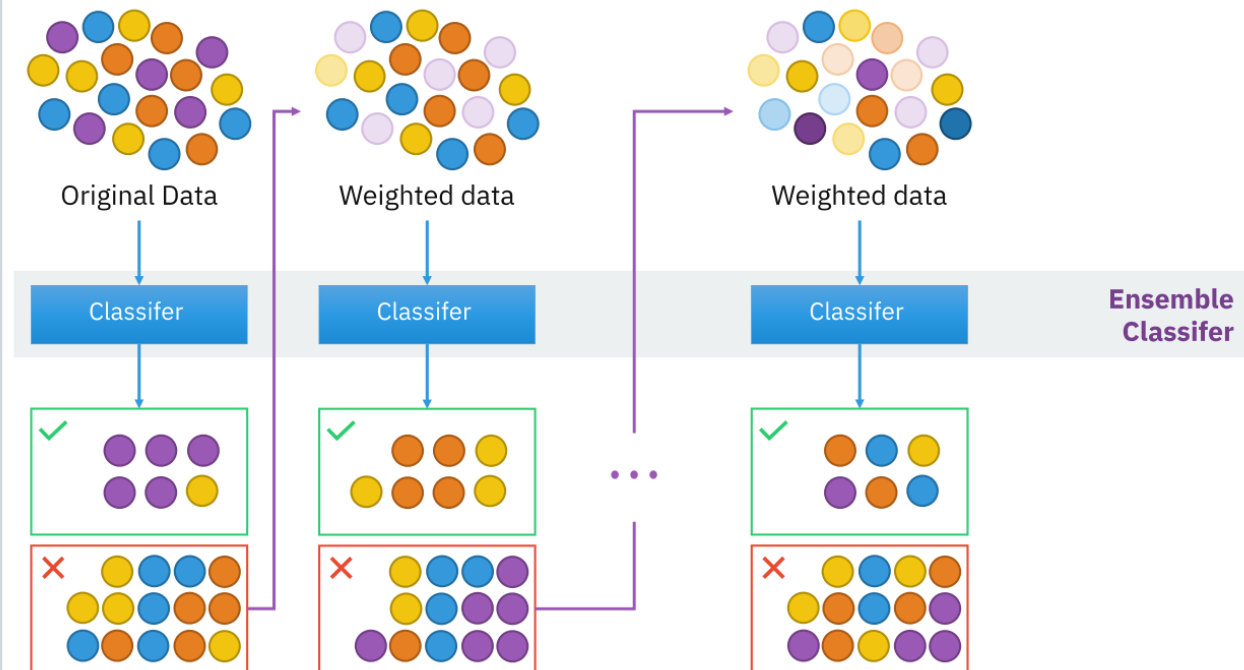
Disadvantages:

- Less interpretable
- More computation



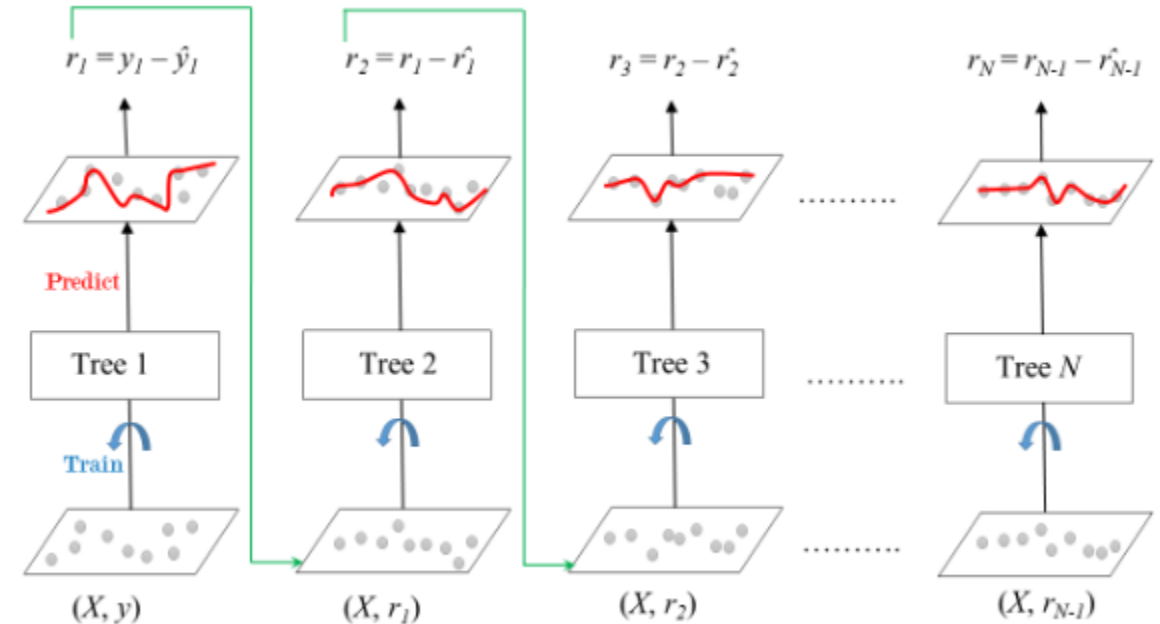
Adaptive Boosting (AdaBoost)

- Boosting: Sequentially trains classifiers
- Weight Adjustment: Emphasizes misclassified instances.
- Final Prediction: weighted sum of the predictions of all weak classifiers.



Gradient Boosting

- Boosting: Sequentially trains classifiers
- Target: the residuals of the previous tree
- Final Prediction: sum of the initial prediction and the contributions of all the weak learners, each multiplied by the learning rate



Wide Range of Boosting Algorithms

Same main idea, different details and implementations:

- Speed and memory usage
- Handling missing values
- Handling categorical values
- Regularization
- Tree pruning
- Parallelism
- Community

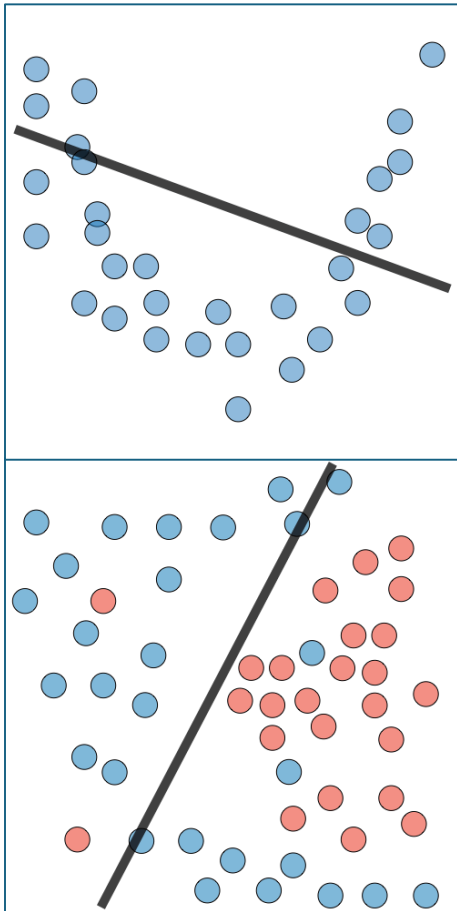
Python packages:

- [sklearn.ensemble](#)
- [CatBoost](#)
- [LightGBM](#)
- [XGBoost](#)

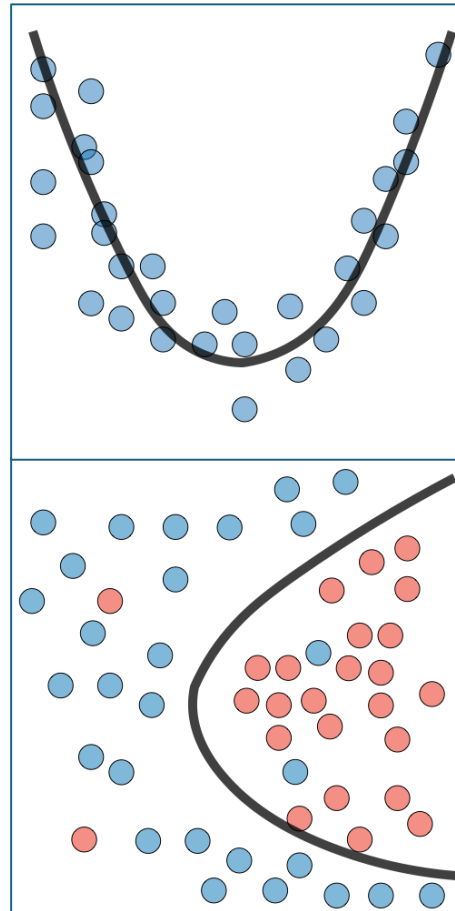
Overfitting and Generalization

Overfitting and Generalization

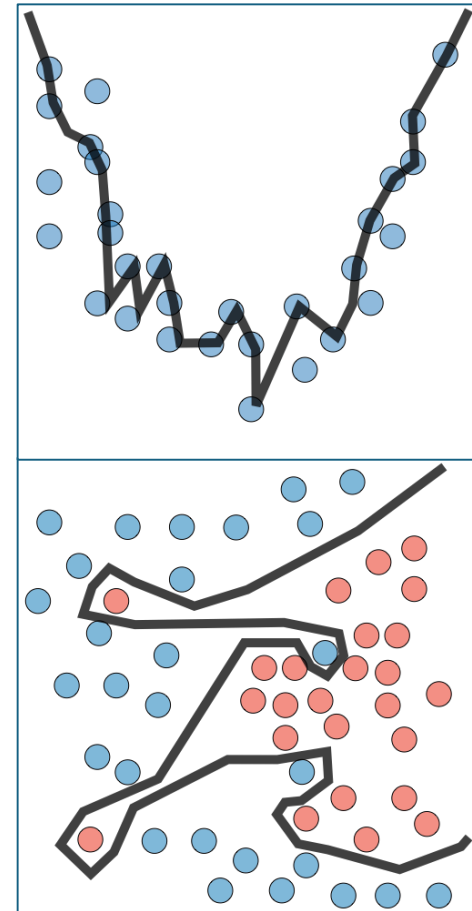
Underfitting



Good Model



Overfitting



Overfitting and Generalization

Sample	Small	Big
Features	Few	Many
Model	Simple	Complex

Possible solutions:

- Collect more data (observations and/or features)
- Enhance data quality (outliers, noise)
- Feature selection, engineering and regularization
- Control model complexity by tuning hyperparameters
- Better evaluation with cross-validation

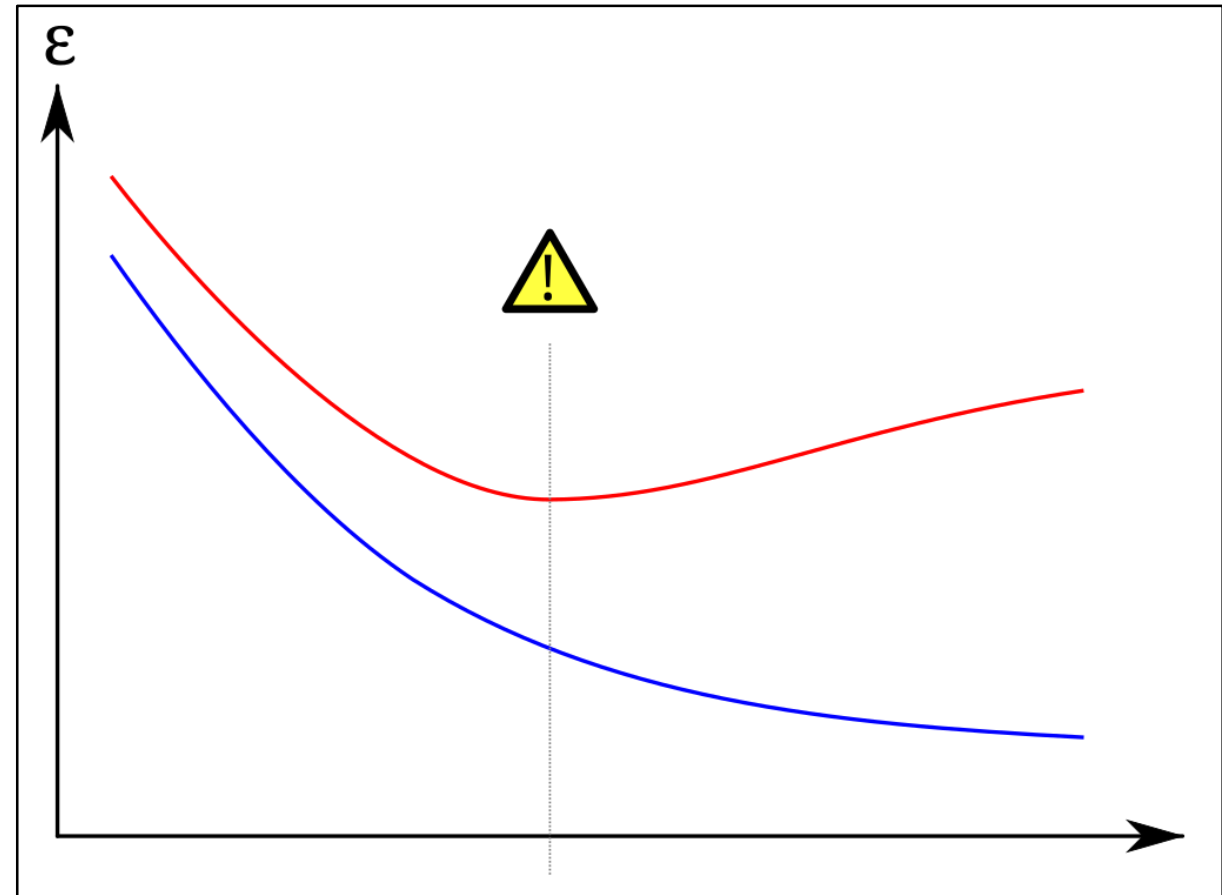
Overfitting and Generalization

Reasons:

- Learning noise
- Memorization
- Sampling issue

Steps:

1. Build a simple model
2. Try to overfit
3. Tune the hyperparameters



Feature Selection

1. Initial feature selection

- Business meaning
- Target leakage
- Predictive power
- Multicollinearity
- Stability

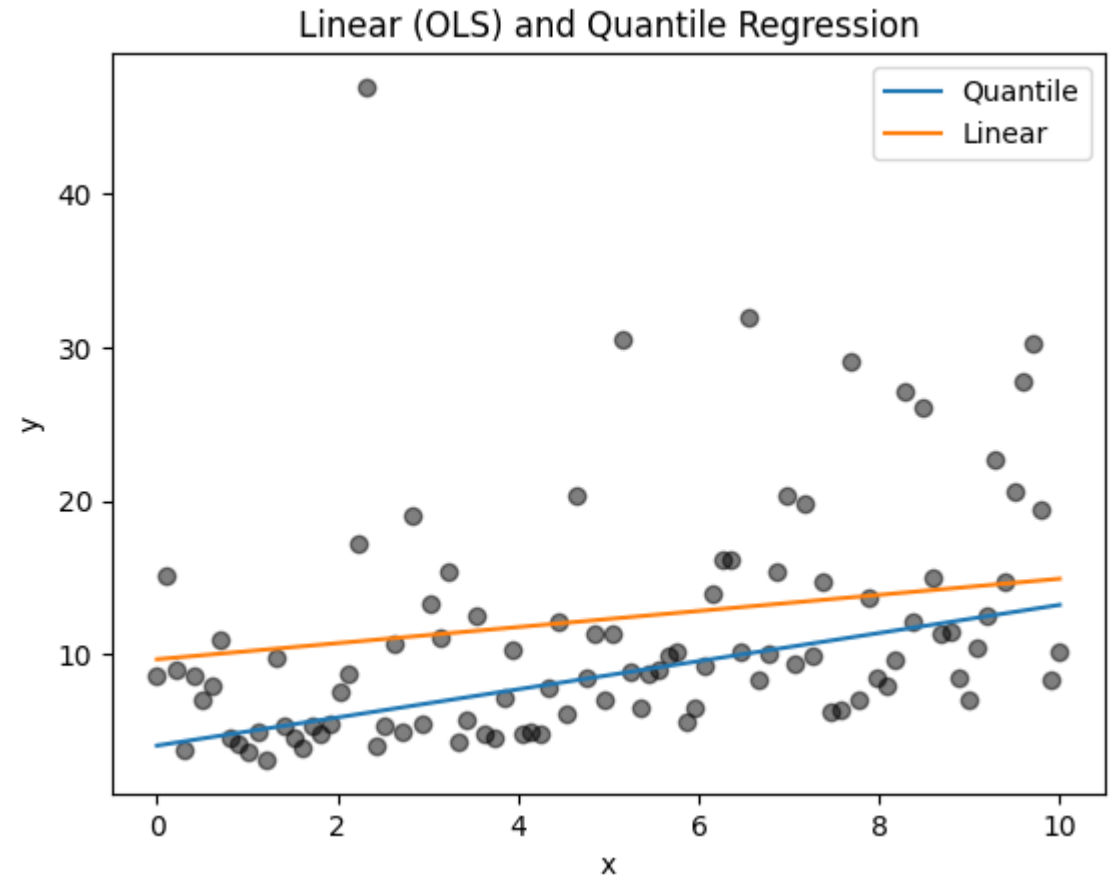
2. Training-time feature selection

- Stepwise feature selection
- L1 Regularization (LASSO regression)
- Feature importances

Quantile Regression

- We can change the loss function from MSE to MAE for some algorithms
- Estimates the median instead of the mean
- Less sensitive to outliers

	RMSE	MAE
OLS Linear Regression	7.2	5.4
Quantile Regression	7.5	4.6



$L_1 - L_2$ Regularization

Lasso (L_1) $Loss_{L_1} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{i=1}^p |\beta_i|$

Ridge (L_2) $Loss_{L_2} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{i=1}^p \beta_i^2$

Elastic net $Loss_{L_1+L_2} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha_1 \sum_{i=1}^p |\beta_i| + \alpha_2 \sum_{i=1}^p \beta_i^2$

- The error on the train sample will be higher with L_1 and L_2 penalties, but we might avoid overfitting and reduce the test error
- Lasso is a technique for feature selection as well
- Feature standardization is beneficial
- The same can be applied for logistic regression, neural networks and XGBoost

Hyperparameter Optimization

Hyperparameter Optimization

Hyperparameters: Parameters that are not learned from the data but set prior to the training process

Hyperparameter Optimization: The process of finding the best set of hyperparameters for a machine learning model to improve its performance.

Importance: The wrong hyperparameters can cause **underfitting** or **overfitting**.

Hyperparameter Optimization

kNN:

- Number of neighbors (k)
- Weights
- Distance metric

SVM:

- C (regularization parameter)
- Kernel
- Kernel parameters (e.g. gamma)

Decision Tree:

- Max depth
- Min samples on leafs
- Max features

Random Forest:

- Tree parameters
- Number of estimators

Gradient Boosting:

- Tree parameters
- Number of estimators
- Learning rate
- Subsample

...and many more

Hyperparameter Optimization

Grid Search: exhaustive search over a specified parameter grid

- Pros: Simple and comprehensive
- Cons: Computationally expensive

Random Search: randomly samples hyperparameters from a specified distribution

- Pros: More efficient than grid search
- Cons: May miss optimal combinations

Bayesian Optimization: uses probabilistic models to predict the performance of hyperparameters.

- Pros: Efficient and can find better hyperparameters.
- Cons: More complex to implement.

Cross-Validation

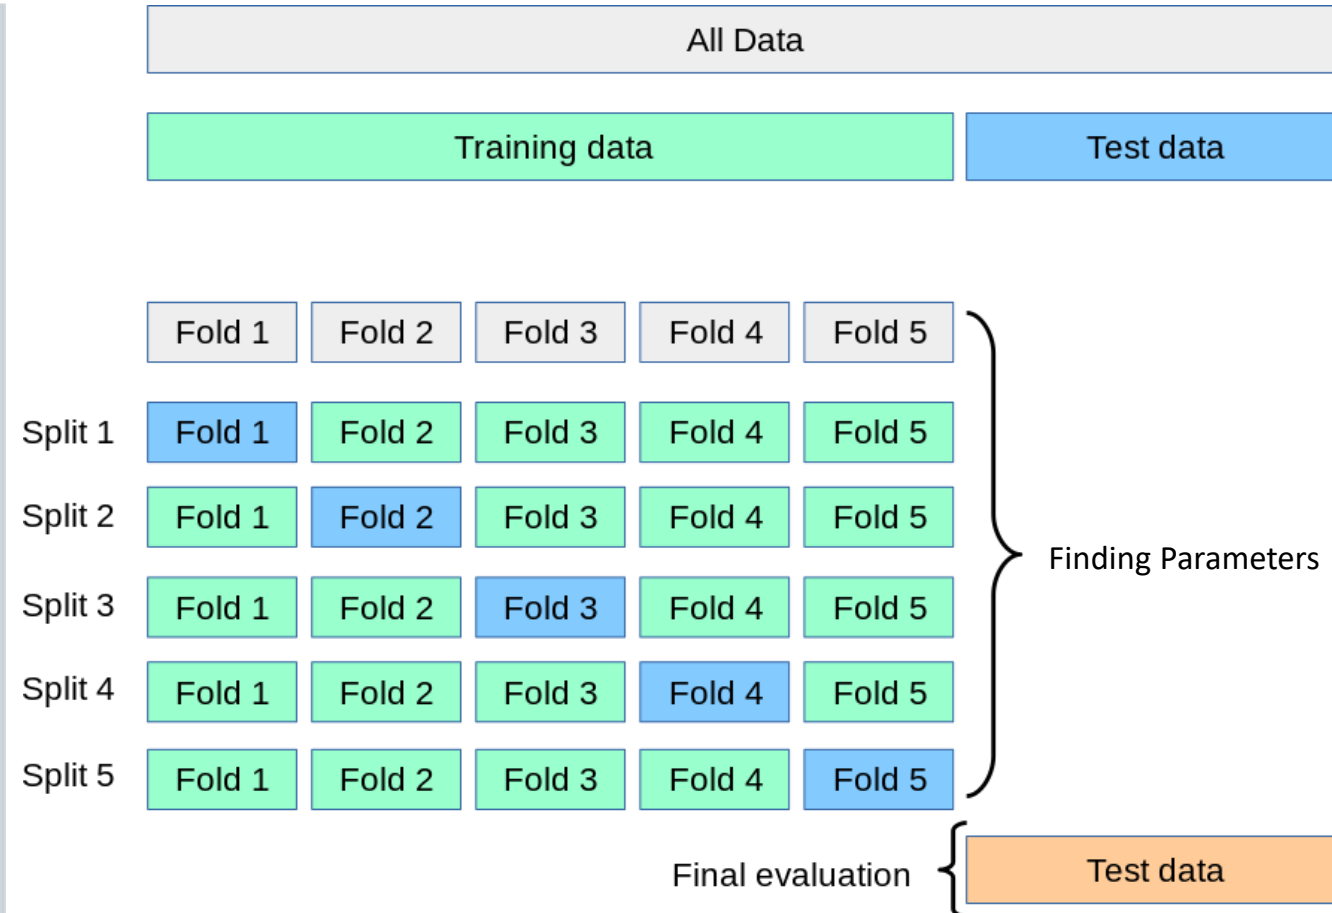
More reliable evaluation, prevents overfitting to the validation sample

K-Fold CV:

- data is divided into k (5-10) equal-sized folds.
- model is trained on k-1 folds and validated on the remaining fold.
- this is repeated k times, with each fold used exactly once as the validation set.

Repeated K-Fold CV: Perform multiple runs of K-Fold CV

Leave-One-Out CV: Each data point is used once as a validation set, and the model is trained on the remaining data



Model Selection – No free lunch theorem

“All models are wrong, but some are useful”

Things to consider:

- Sample size, number of features, problem complexity
- Training/Inference runtime, memory, compute
- Robustness, extrapolation capabilities
- Interpretability

Start simple. Experiment. Use domain knowledge.

Overfitting and Model Tuning Summary

1. Sufficient Sample Size
2. Data Completeness (having the right features)
3. Data Quality (outliers, noise)
4. Feature Selection (no irrelevant features)
5. Feature Engineering (helping the model with domain knowledge)
6. Model Selection (according to the data and the problem)
7. Loss function and evaluation metric according to the business problem
8. Regularization
9. Hyperparameter Search
10. Reliable Evaluation with Cross-Validation

Thank you for your attention!

Your feedback would be much appreciated:



Any Questions?



Gergely Zsombor Haász



haasz.zsombi@gmail.com